

Time Series Analysis: Cambodia Rainfall Forecasting

Course: Time Series Analysis

Group: 01

Instructor: Dr. SIM Tepmomy

Date: January 2026

Corrected Version - All Modern APIs

This notebook uses:

-  Modern `statsmodels.tsa.arima.model.ARIMA` and `SARIMAX` (no deprecated `disp` parameter)
 -  Proper train-test split (no data leakage)
 -  Out-of-sample rolling forecasts for evaluation
 -  Seasonal decomposition and seasonal differencing
 -  Residual diagnostics (ACF, Ljung-Box, normality tests)
 -  Model comparison with proper metrics
-

Table of Contents

1. Import Libraries
 2. Data Loading and Exploration
 3. Data Preprocessing and Monthly Aggregation
 4. Exploratory Data Analysis (EDA)
 5. Time Series Decomposition
 6. Stationarity Testing (ADF)
 7. Seasonal Differencing
 8. ACF and PACF Analysis
 9. Train-Test Split (No Leakage)
 10. Model Fitting with Out-of-Sample Evaluation
 - AR Models
 - MA Models
 - ARMA Models
 - ARIMA Models
 - SARIMA Models
 11. Model Comparison
 12. Residual Diagnostics
 13. Final Forecasting
 14. Conclusions
-

1. Import Libraries

All modern APIs - no deprecated functions.

```
In [ ]: import warnings

# Data manipulation
import numpy as np
import pandas as pd

# Visualization
import matplotlib.pyplot as plt
import seaborn as sns

plt.style.use("seaborn-v0_8-darkgrid")
sns.set_palette("husl")

# Time series analysis - MODERN APIs only
from statsmodels.tsa.stattools import adfuller, acf, pacf
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
from statsmodels.tsa.seasonal import seasonal_decompose
from statsmodels.tsa.arima.model import ARIMA # Modern ARIMA
from statsmodels.tsa.statespace.sarimax import SARIMAX # Modern SARIMA
from statsmodels.stats.diagnostic import acorr_ljungbox

# >>> add these imports <<<
from statsmodels.tools.sm_exceptions import ValueWarning, ConvergenceWarning

# Machine Learning metrics
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

# Set random seed
np.random.seed(42)

# Suppress specific statsmodels warnings
warnings.filterwarnings("ignore", category=ValueWarning)
warnings.filterwarnings("ignore", category=ConvergenceWarning)

import warnings
import statsmodels

warnings.filterwarnings("ignore")

# Completely silence all UserWarning coming from statsmodels
warnings.filterwarnings(
    "ignore", category=UserWarning, module=r"statsmodels\tsa\statespace\sarimax"
)
```

2. Data Loading and Initial Exploration

Mathematical Background: Time Series Definition

A **time series** is a sequence of observations y_t indexed by time t , where $t = 1, 2, \dots, T$.

$$y_t = f(t) + \varepsilon_t$$

where:

- $f(t)$ represents the systematic component (trend + seasonality)
- ε_t represents random noise (white noise)

Key Properties:

1. Mean: $\mu_t = \mathbb{E}[y_t]$
2. Variance: $\sigma_t^2 = \text{Var}(y_t)$
3. Autocovariance: $\gamma(s, t) = \text{Cov}(y_s, y_t)$

```
In [2]: # Load Cambodia rainfall dataset from HDX
url = "https://data.humdata.org/dataset/8fa90d2b-a88e-414d-84a1-50d6bc773542/resource/67c4f3d6-f600-4699-9a67-0de20d6a1b0b/download/khm-rainfall-subnat-full.csv"

print("Loading Cambodia rainfall dataset...")
df_raw = pd.read_csv(url)
print(f"✓ Dataset loaded successfully!")
print(f"Shape: {df_raw.shape}")
print(f"\nFirst few rows:")
df_raw.head()
```

Loading Cambodia rainfall dataset...
✓ Dataset loaded successfully!
Shape: (358241, 15)

First few rows:

```
Out[2]:   date  adm_level  adm_id  PCODE  n_pixels  rfh  rfh_avg  r1h  r1h_avg  r3h  r3h_avg  rfq  r1q  r3q  version
0  1981-01-01      1  900411    KH15       2.0    1.0   2.400000  NaN  NaN  NaN  81.081080  NaN  NaN  final
1  1981-01-11      1  900411    KH15       2.0    0.5   1.000000  NaN  NaN  NaN  91.666670  NaN  NaN  final
2  1981-01-21      1  900411    KH15       2.0    3.0   2.250000  4.5  5.650000  NaN  110.344826  89.20188  NaN  final
3  1981-02-01      1  900411    KH15       2.0    4.5   2.483333  8.0  5.733333  NaN  126.948780  121.11801  NaN  final
4  1981-02-11      1  900411    KH15       2.0    4.5   4.150000  12.0  8.883333  NaN  103.825140  122.44898  NaN  final
```

```
In [3]: # Dataset information
print("=" * 60)
print("Dataset Information")
print("=" * 60)
df_raw.info()
print("\n" + "=" * 60)
print("Summary Statistics")
print("=" * 60)
df_raw.describe()
```

```
=====
Dataset Information
=====
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 358241 entries, 0 to 358240
Data columns (total 15 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   date        358241 non-null   object 
 1   adm_level   358241 non-null   int64  
 2   adm_id      358241 non-null   int64  
 3   PCODE       358241 non-null   object 
 4   n_pixels    358241 non-null   float64
 5   rfh         358241 non-null   float64
 6   rfh_avg     358241 non-null   float64
 7   r1h         357799 non-null   float64
 8   r1h_avg     357799 non-null   float64
 9   r3h         356473 non-null   float64
 10  r3h_avg     356473 non-null   float64
 11  rfq         358241 non-null   float64
 12  r1q         357799 non-null   float64
 13  r3q         356473 non-null   float64
 14  version     358241 non-null   object 
dtypes: float64(10), int64(2), object(3)
memory usage: 41.0+ MB
```

```
=====
Summary Statistics
=====
```

	adm_level	adm_id	n_pixels	rfh	rfh_avg	r1h	r1h_avg	r3h	r3h_avg	rfq	r1q	r3q
count	358241.000000	3.582410e+05	358241.000000	358241.000000	358241.000000	357799.000000	357799.000000	356473.000000	356473.000000	358241.000000	357799.000000	356473.000000
mean	1.882353	9.938594e+05	53.918552	50.533208	50.446255	151.781830	151.475458	456.628298	455.553361	99.503847	99.195325	99.530930
std	0.322190	3.411807e+04	87.656551	53.591433	44.794200	142.794801	130.781864	369.396604	352.300300	48.516597	39.896489	31.625497
min	1.000000	9.004110e+05	1.000000	0.000000	0.000000	0.000000	0.019672	3.000000	14.302381	7.340091	6.648260	11.726528
25%	2.000000	1.006248e+06	11.000000	6.000000	11.220963	25.300001	35.766666	137.488890	148.483340	68.956610	74.300772	81.343100
50%	2.000000	1.006303e+06	22.000000	37.600000	43.526670	130.000000	137.500000	420.333300	427.566680	90.942500	94.072740	97.329310
75%	2.000000	1.006361e+06	50.000000	77.000000	77.346940	231.333320	229.926350	658.777800	644.513300	119.939100	116.974300	113.842650
max	2.000000	1.006416e+06	467.000000	803.608700	356.183320	1743.347900	941.975000	3548.782700	2363.883300	816.339500	511.743840	352.051330

3. Data Preprocessing and Monthly Aggregation

Mathematical Background: Temporal Aggregation

When aggregating from daily/dekadal to monthly data:

$$Y * m = \sum *d \in my_d$$

where Y_m is total rainfall in month m , and y_d is rainfall on day/dekad d .

Properties:

- Reduces noise (smoothing effect)
- Captures longer-term patterns
- More suitable for seasonal modeling

```
In [4]: # Convert date to datetime
df_raw["date"] = pd.to_datetime(df_raw["date"])

# Focus on provincial level (adm_Level = 1)
df_province = df_raw[df_raw["adm_level"] == 1].copy()

print(f"Provincial-level records: {len(df_province)}")
print(f"Number of provinces: {df_province['PCODE'].nunique()}")
print(f"Provinces: {sorted(df_province['PCODE'].unique())}")
```

```
Provincial-level records: 42146
Number of provinces: 25
Provinces: ['KH01', 'KH02', 'KH03', 'KH04', 'KH05', 'KH06', 'KH07', 'KH08', 'KH09', 'KH10', 'KH11', 'KH12', 'KH13', 'KH14', 'KH15', 'KH16', 'KH17', 'KH18', 'KH19', 'KH20', 'KH21', 'KH22', 'KH23', 'KH24', 'KH25']
```

```
In [5]: # Select province with most complete data
province_counts = df_province.groupby("PCODE").size()
selected_province = province_counts.idxmax()

print(f"Selected province: {selected_province}")
print(f"Records: {province_counts[selected_province]}")

df_selected = df_province[df_province["PCODE"] == selected_province].copy()
df_selected = df_selected.sort_values("date").reset_index(drop=True)
df_selected = df_selected[["date", "rfh"]].copy()
df_selected.columns = ["date", "rainfall"]

print(f"\nData range: {df_selected['date'].min()} to {df_selected['date'].max()}")
df_selected.head()
```

```
Selected province: KH15
Records: 3242
```

```
Data range: 1981-01-01 00:00:00 to 2026-01-01 00:00:00
```

```
Out[5]:
```

	date	rainfall
0	1981-01-01	1.000000
1	1981-01-01	1.393782
2	1981-01-11	0.500000
3	1981-01-11	1.489637
4	1981-01-21	3.000000

```
In [6]: # Handle missing values - modern pandas syntax
print("=" * 60)
print("Missing Value Analysis")
print("=" * 60)
print(f"Missing values: {df_selected['rainfall'].isna().sum()}")
print(
    f"Percentage: {df_selected['rainfall'].isna().sum() / len(df_selected) * 100:.2f}%"
```

```
)  
  
# Use modern pandas methods (no deprecated 'method' parameter)  
df_selected["rainfall"] = df_selected["rainfall"].ffill().bfill()  
  
print(f"After imputation: {df_selected['rainfall'].isna().sum()}")
```

```
=====  
Missing Value Analysis  
=====
```

```
Missing values: 0  
Percentage: 0.00%  
After imputation: 0
```

```
In [7]: # Aggregate to monthly  
df_selected["year_month"] = df_selected["date"].dt.to_period("M")  
  
df_monthly = df_selected.groupby("year_month").agg({"rainfall": "sum"}).reset_index()  
  
df_monthly["date"] = df_monthly["year_month"].dt.to_timestamp()  
df_monthly = df_monthly[["date", "rainfall"]].copy()  
  
print("=" * 60)  
print("Monthly Data")  
print("=" * 60)  
print(f"Total months: {len(df_monthly)}")  
print(f"Range: {df_monthly['date'].min()} to {df_monthly['date'].max()}")  
print(f"\nFirst rows:")  
df_monthly.head(10)
```

```
=====  
Monthly Data  
=====
```

```
Total months: 541  
Range: 1981-01-01 00:00:00 to 2026-01-01 00:00:00
```

```
First rows:
```

```
Out[7]:
```

	date	rainfall
0	1981-01-01	17.272021
1	1981-02-01	76.948185
2	1981-03-01	53.823835
3	1981-04-01	149.800519
4	1981-05-01	333.725390
5	1981-06-01	569.883415
6	1981-07-01	589.834190
7	1981-08-01	582.481870
8	1981-09-01	486.427453
9	1981-10-01	472.746120

```
In [8]: # Create time series  
ts_data = df_monthly.set_index("date")["rainfall"]  
ts_data = ts_data.asfreq("MS")
```

```

print(f"\nTime series shape: {ts_data.shape}")
print(f"Frequency: {ts_data.index.freq}")
ts_data.head()

Time series shape: (541,)
Frequency: <MonthBegin>

Out[8]: date
1981-01-01    17.272021
1981-02-01    76.948185
1981-03-01    53.823835
1981-04-01   149.800519
1981-05-01   333.725390
Freq: MS, Name: rainfall, dtype: float64

```

4. Exploratory Data Analysis (EDA)

Mathematical Background: Descriptive Statistics

For time series $\{y_t\} * t = 1^T$:

Sample Mean:

$$\bar{y} = \frac{1}{T} \sum_t y_t = 1^T y_t$$

Sample Variance:

$$s^2 = \frac{1}{T-1} \sum_t (y_t - \bar{y})^2$$

Coefficient of Variation:

$$CV = \frac{s}{\bar{y}} \times 100\%$$

Skewness:

$$\text{Skew} = \frac{1}{T} \sum_t \left(\frac{y_t - \bar{y}}{s} \right)^3$$

Kurtosis:

$$\text{Kurt} = \frac{1}{T} \sum_t \left(\frac{y_t - \bar{y}}{s} \right)^4$$

```

In [9]: # Descriptive statistics
print("=" * 60)
print("Descriptive Statistics")
print("=" * 60)
stats = {
    "Count": len(ts_data),
    "Mean": ts_data.mean(),
    "Median": ts_data.median(),
}

```

```

        "Std Dev": ts_data.std(),
        "Variance": ts_data.var(),
        "Min": ts_data.min(),
        "Max": ts_data.max(),
        "Range": ts_data.max() - ts_data.min(),
        "Q1": ts_data.quantile(0.25),
        "Q3": ts_data.quantile(0.75),
        "IQR": ts_data.quantile(0.75) - ts_data.quantile(0.25),
        "CV (%)": (ts_data.std() / ts_data.mean()) * 100,
        "Skewness": ts_data.skew(),
        "Kurtosis": ts_data.kurtosis(),
    }

    for key, value in stats.items():
        print(f"{key}: {value:.4f}")

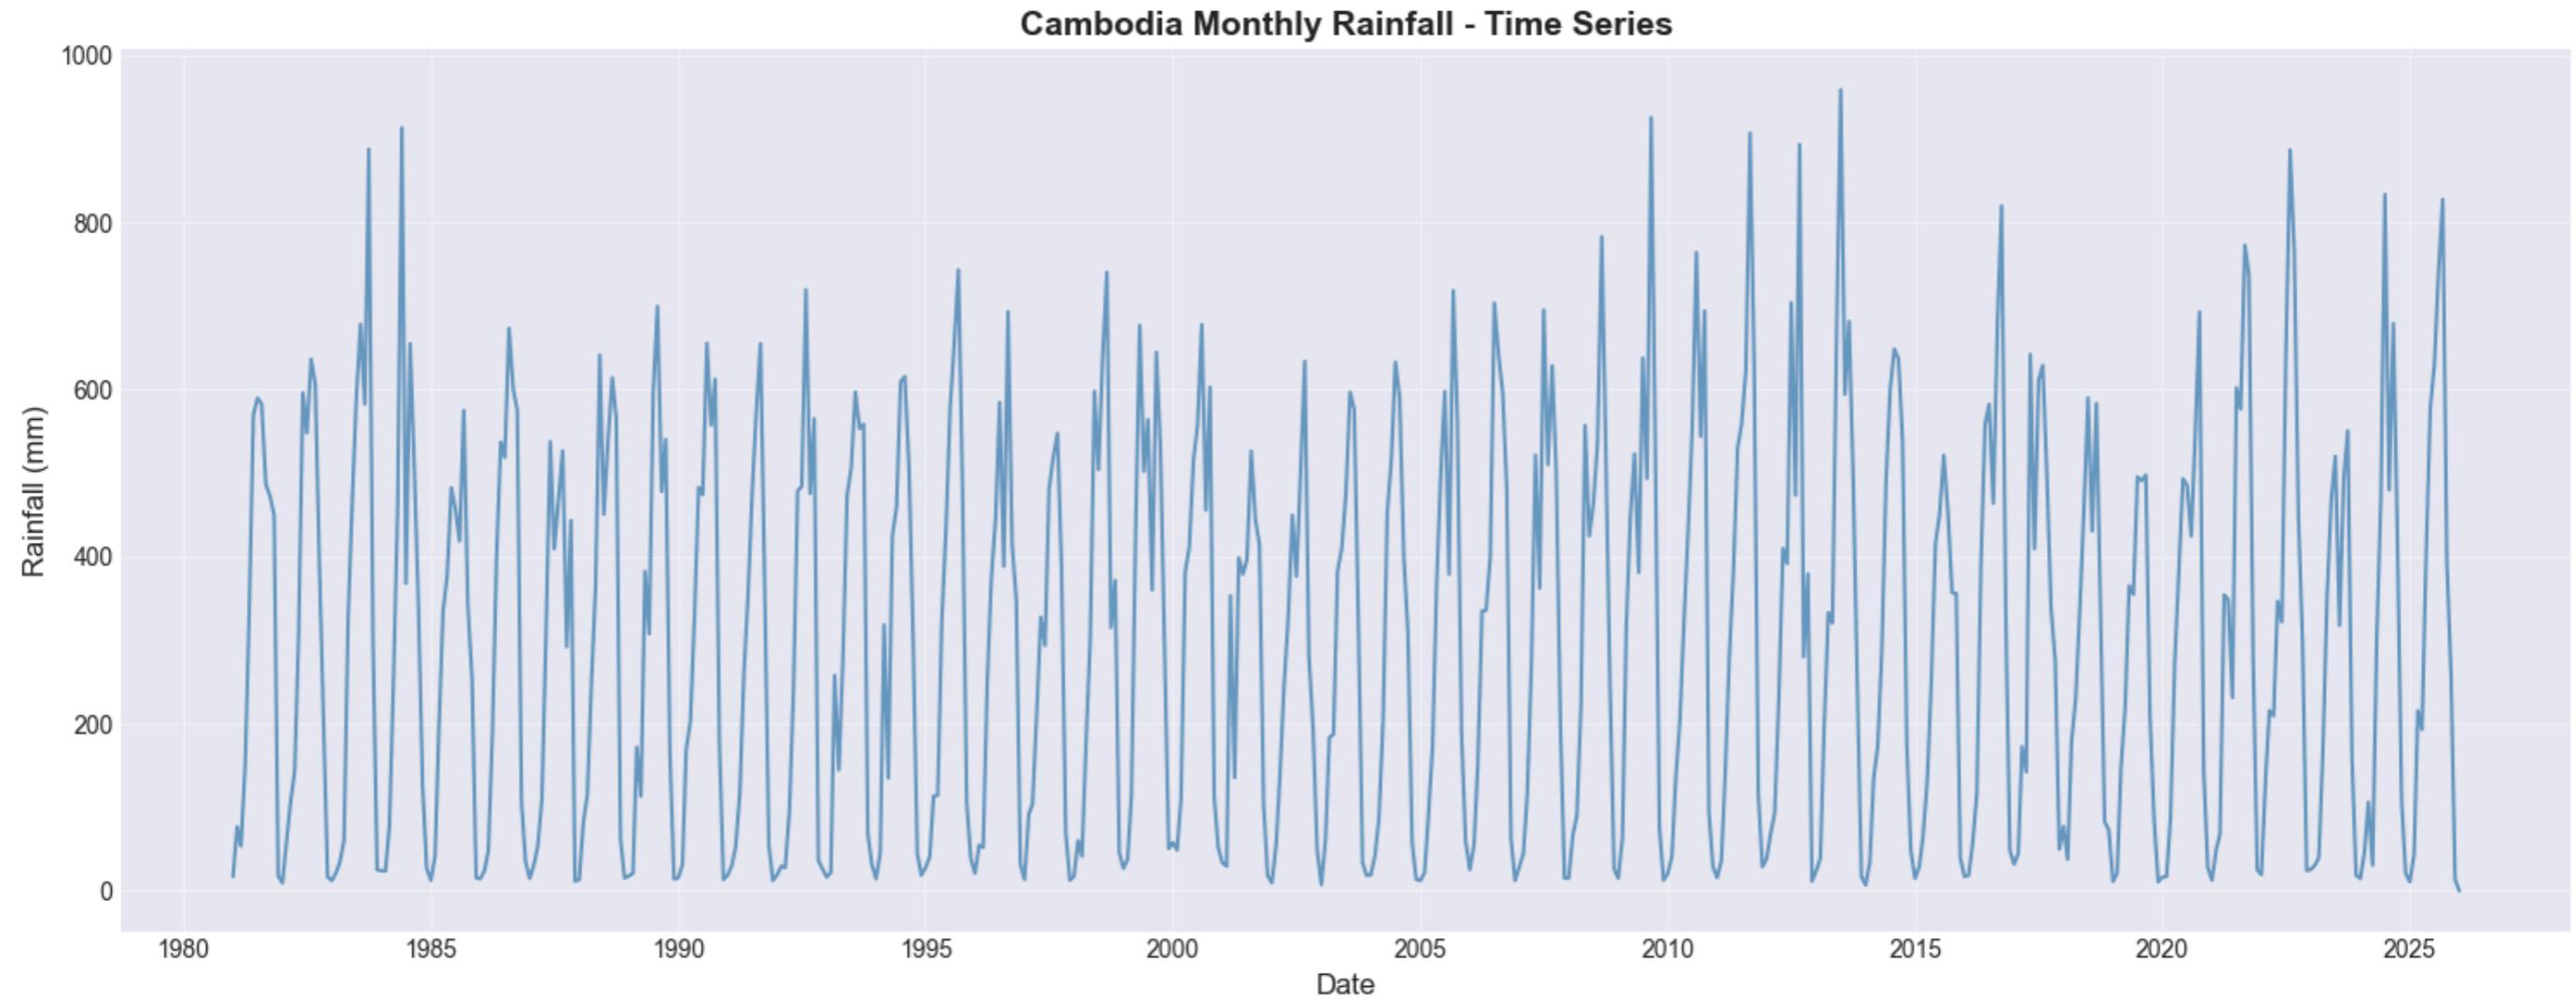
```

```
=====
Descriptive Statistics
=====
Count      : 541.0000
Mean       : 306.2502
Median     : 308.3109
Std Dev    : 239.4916
Variance   : 57356.2467
Min        : 0.0104
Max        : 958.9870
Range      : 958.9767
Q1         : 55.2539
Q3         : 501.9611
IQR        : 446.7072
CV (%)     : 78.2013
Skewness   : 0.3308
Kurtosis   : -1.0155
```

```
In [10]: # Time series plot
fig, ax = plt.subplots(figsize=(15, 6))

ax.plot(ts_data.index, ts_data.values, linewidth=1.5, color="steelblue", alpha=0.8)
ax.set_title("Cambodia Monthly Rainfall - Time Series", fontsize=14, fontweight="bold")
ax.set_xlabel("Date", fontsize=12)
ax.set_ylabel("Rainfall (mm)", fontsize=12)
ax.grid(True, alpha=0.3)

plt.tight_layout()
plt.savefig("timeseries_plot.png", dpi=300, bbox_inches="tight")
plt.show()
print("✓ Saved: timeseries_plot.png")
```



✓ Saved: timeseries_plot.png

5. Time Series Decomposition

Mathematical Background: Additive Decomposition

$$y_t = T_t + S_t + R_t$$

where:

- T_t = Trend component
- S_t = Seasonal component
- R_t = Residual (irregular) component

Seasonal Period: $s = 12$ months (annual seasonality)

```
In [11]: # Seasonal decomposition
decomposition = seasonal_decompose(ts_data, model="additive", period=12)

fig, axes = plt.subplots(4, 1, figsize=(15, 12))

# Original
axes[0].plot(ts_data, linewidth=1.5, color="black")
```

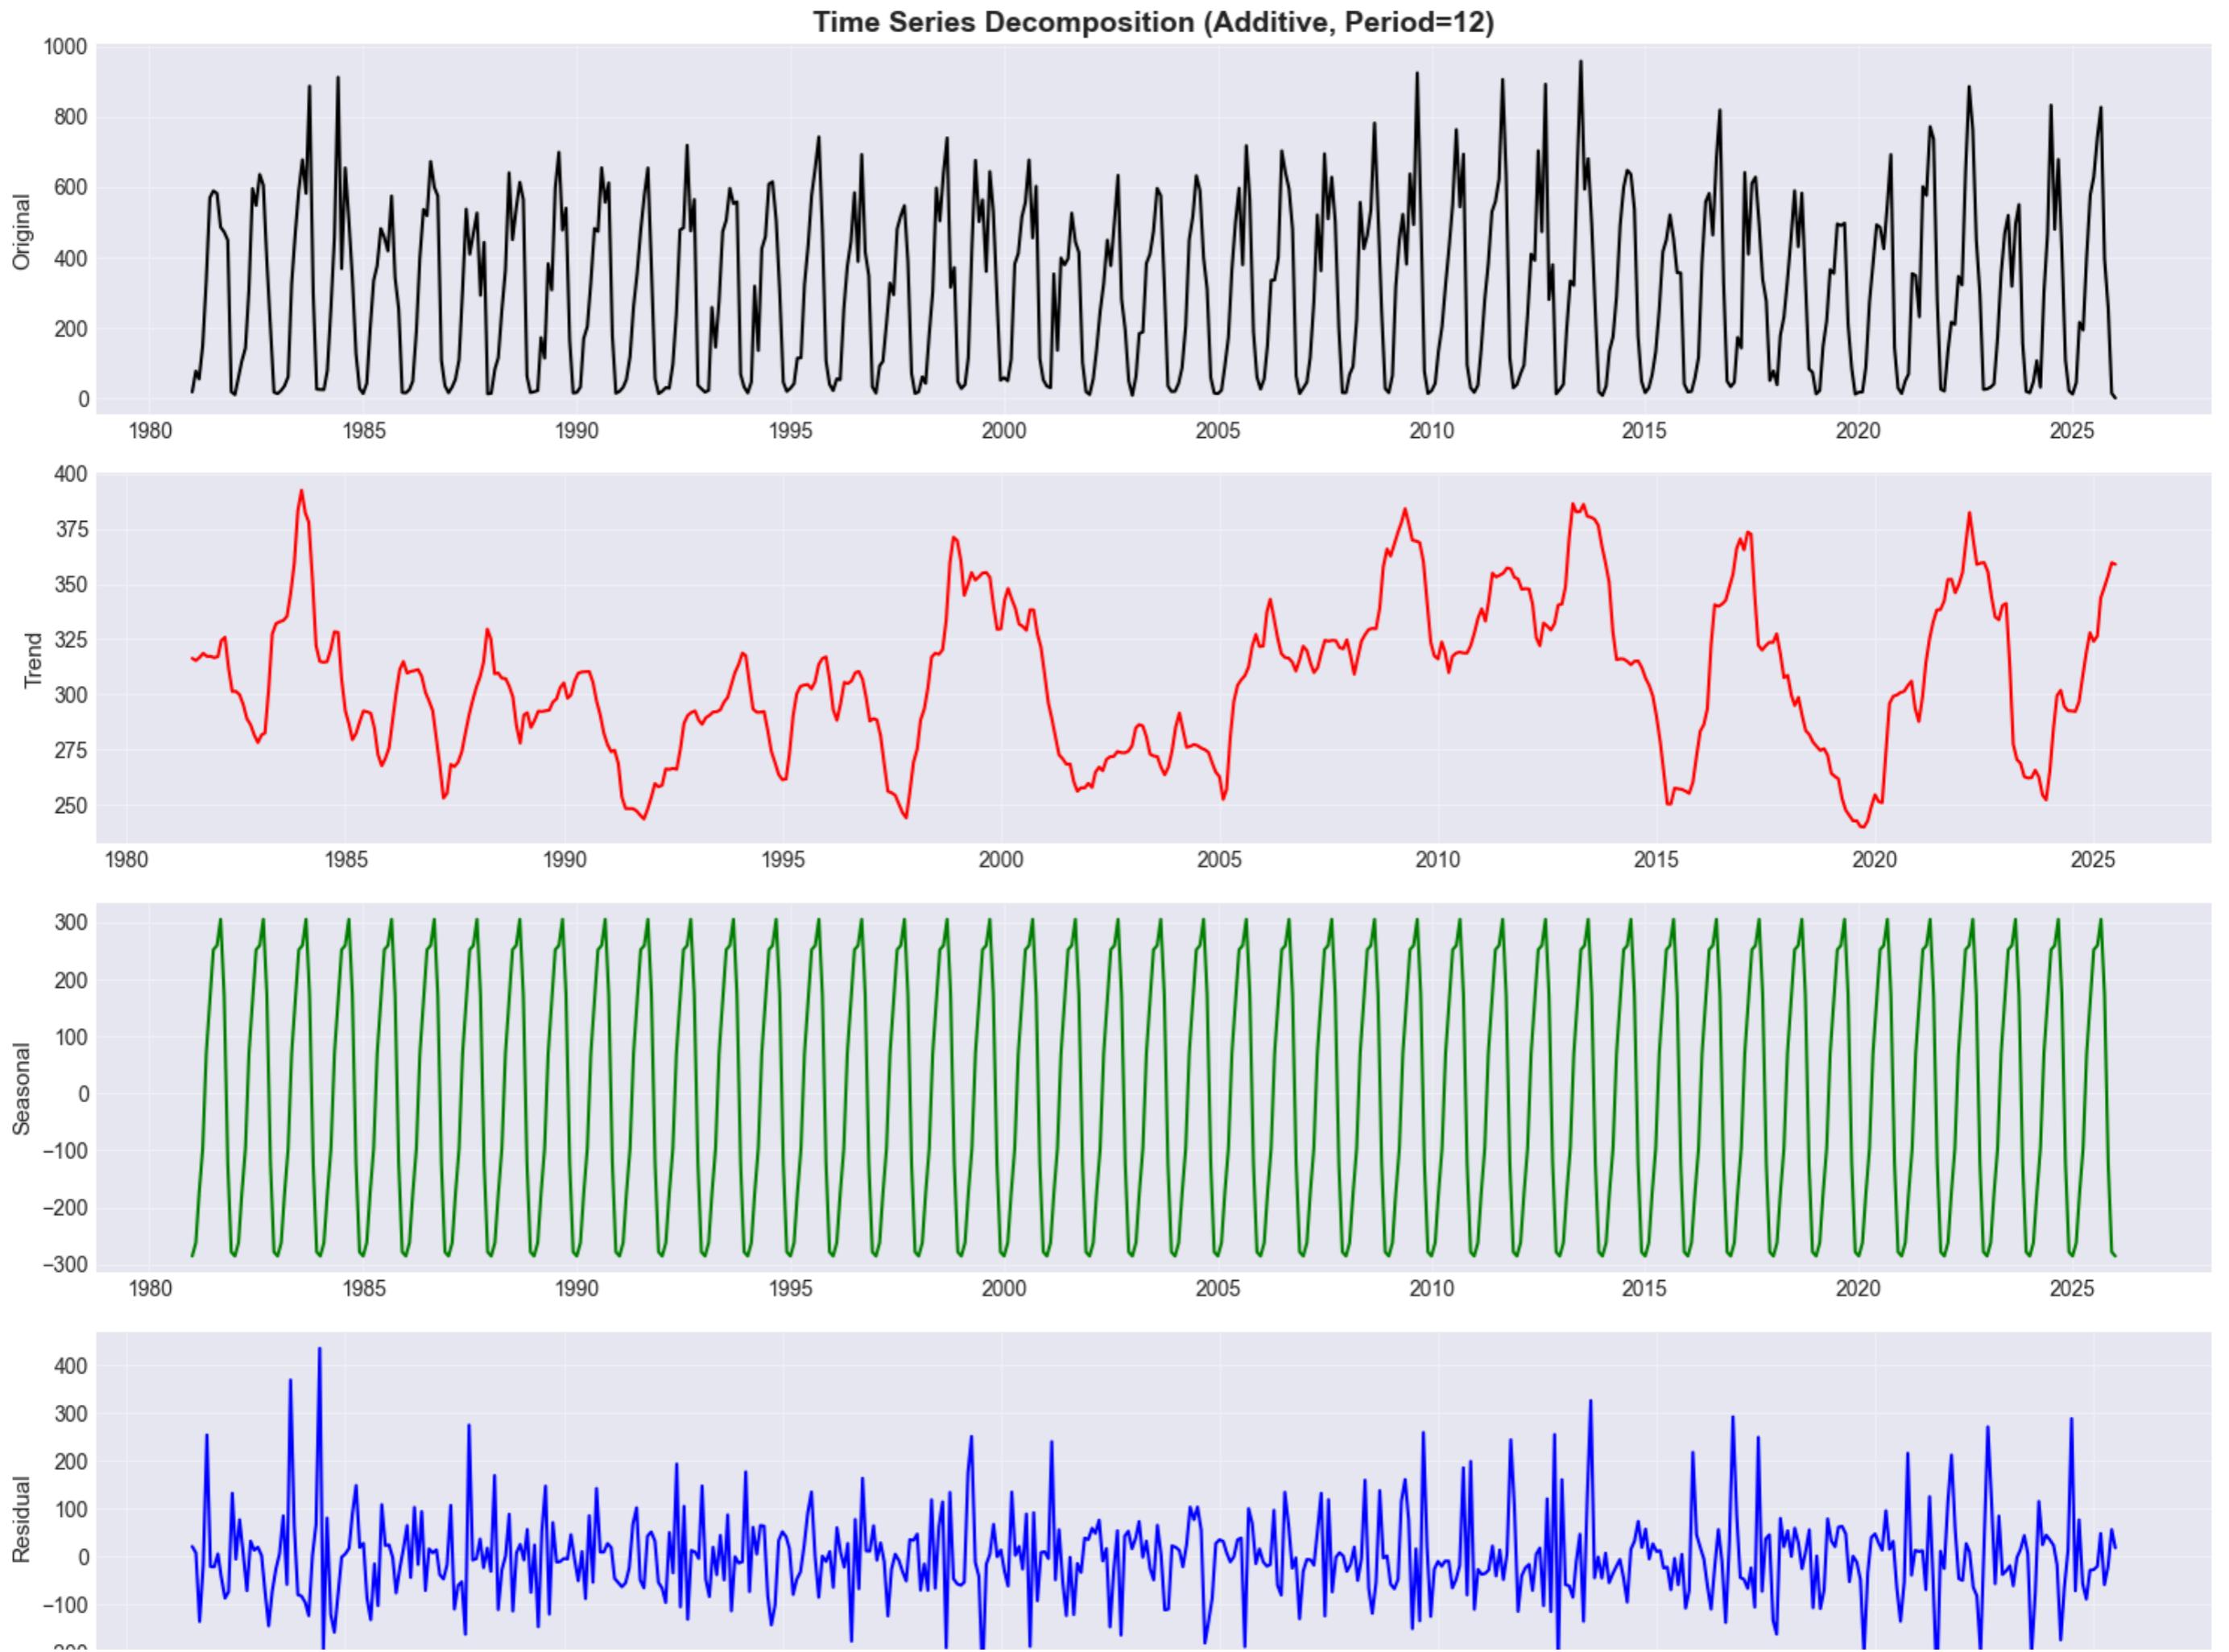
```
axes[0].set_ylabel("Original", fontsize=11)
axes[0].set_title(
    "Time Series Decomposition (Additive, Period=12)", fontsize=14, fontweight="bold"
)
axes[0].grid(True, alpha=0.3)

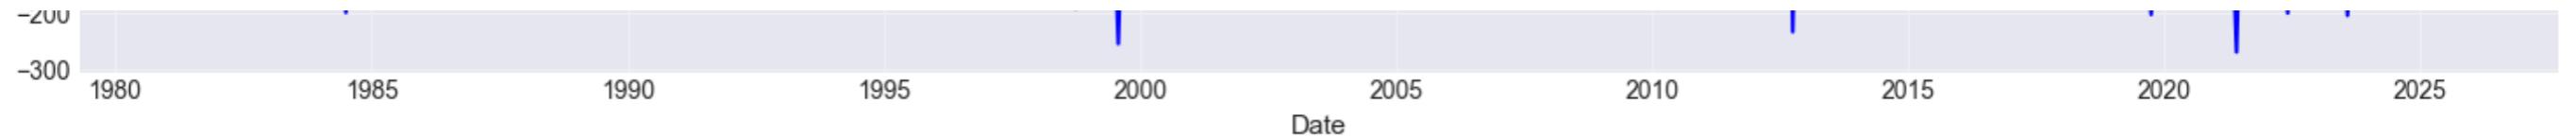
# Trend
axes[1].plot(decomposition.trend, linewidth=1.5, color="red")
axes[1].set_ylabel("Trend", fontsize=11)
axes[1].grid(True, alpha=0.3)

# Seasonal
axes[2].plot(decomposition.seasonal, linewidth=1.5, color="green")
axes[2].set_ylabel("Seasonal", fontsize=11)
axes[2].grid(True, alpha=0.3)

# Residual
axes[3].plot(decomposition.resid, linewidth=1.5, color="blue")
axes[3].set_ylabel("Residual", fontsize=11)
axes[3].set_xlabel("Date", fontsize=11)
axes[3].grid(True, alpha=0.3)

plt.tight_layout()
plt.savefig("decomposition.png", dpi=300, bbox_inches="tight")
plt.show()
print("✓ Saved: decomposition.png")
```





✓ Saved: decomposition.png

6. Stationarity Testing (ADF Test)

Mathematical Background: Augmented Dickey-Fuller Test

Null Hypothesis (H_0): The series has a unit root (non-stationary)

Alternative Hypothesis (H_1): The series is stationary

Test Statistic:

$$\Delta y * t = \alpha + \beta t + \gamma y * t - 1 + \sum *i = 1^p \delta_i \Delta y * t - i + \varepsilon_t$$

Decision Rule:

- If p-value < 0.05: Reject H_0 (series is stationary)
- If p-value ≥ 0.05 : Fail to reject H_0 (series is non-stationary)

```
In [12]: def adf_test(series, name="Series"):
    """Perform Augmented Dickey-Fuller test"""
    result = adfuller(series.dropna(), autolag="AIC")

    print(f"\n{'='*60}")
    print(f"ADF Test Results: {name}")
    print(f"{'='*60}")
    print(f"ADF Statistic: {result[0]:.6f}")
    print(f"p-value: {result[1]:.6f}")
    print(f"# Lags Used: {result[2]}")
    print(f"# Observations: {result[3]}")

    print(f"\nCritical Values:")
    for key, value in result[4].items():
        print(f"  {key}: {value:.3f}")

    if result[1] < 0.05:
        print(
            f"\n/\ Conclusion: Series is STATIONARY (p-value = {result[1]:.6f} < 0.05)"
        )
    else:
        print(
            f"\nX Conclusion: Series is NON-STATIONARY (p-value = {result[1]:.6f} \geq 0.05)"
        )

    return result

# Test original series
adf_result_original = adf_test(ts_data, "Original Series")
```

```
=====
ADF Test Results: Original Series
=====
ADF Statistic:      -5.030478
p-value:           0.000019
# Lags Used:      17
# Observations:   523

Critical Values:
 1%   : -3.443
 5%   : -2.867
10%  : -2.570

✓ Conclusion: Series is STATIONARY (p-value = 0.000019 < 0.05)
```

7. Seasonal Differencing

Mathematical Background: Seasonal Difference

$$\nabla^s y_t = y_t - y_{t-s}$$

where $s = 12$ (seasonal period)

This removes seasonal patterns and helps achieve stationarity.

```
In [13]: # Seasonal differencing (lag = 12)
ts_seasonal_diff = ts_data.diff(12)

# Test stationarity after seasonal differencing
adf_result_seasonal = adf_test(
    ts_seasonal_diff, "Seasonally Differenced Series (lag=12)"
)

=====

ADF Test Results: Seasonally Differenced Series (lag=12)
=====
ADF Statistic:      -9.726492
p-value:           0.000000
# Lags Used:      11
# Observations:   517

Critical Values:
 1%   : -3.443
 5%   : -2.867
10%  : -2.570

✓ Conclusion: Series is STATIONARY (p-value = 0.000000 < 0.05)
```

```
In [14]: # Plot seasonally differenced series
fig, ax = plt.subplots(figsize=(15, 6))

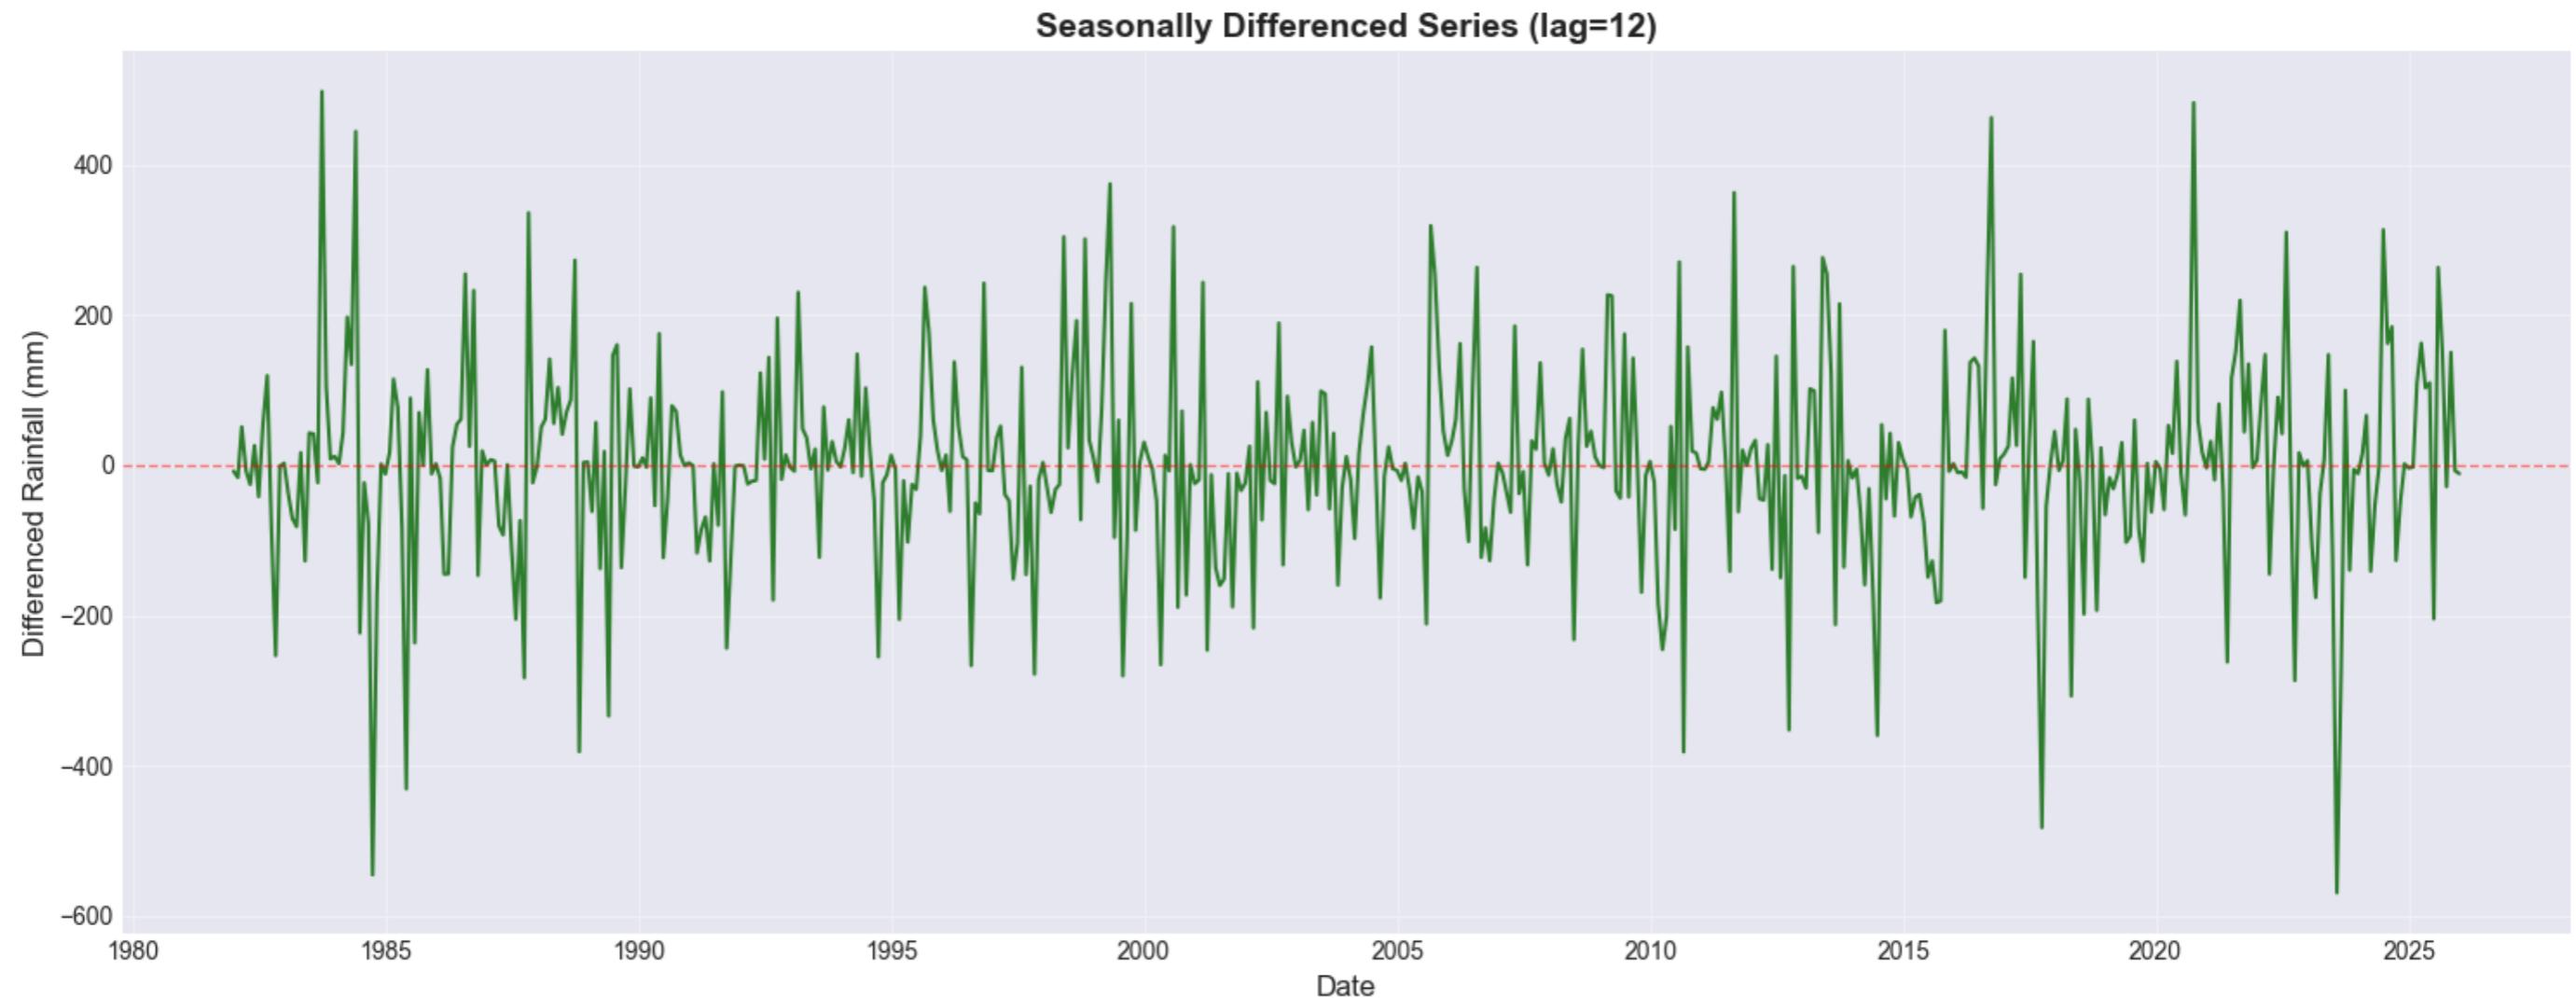
ax.plot(ts_seasonal_diff, linewidth=1.5, color="darkgreen", alpha=0.8)
ax.axhline(y=0, color="red", linestyle="--", linewidth=1, alpha=0.5)
ax.set_title("Seasonally Differenced Series (lag=12)", fontsize=14, fontweight="bold")
ax.set_xlabel("Date", fontsize=12)
ax.set_ylabel("Differenced Rainfall (mm)", fontsize=12)
```

```

ax.grid(True, alpha=0.3)

plt.tight_layout()
plt.savefig("seasonal_differencing.png", dpi=300, bbox_inches="tight")
plt.show()
print("✓ Saved: seasonal_differencing.png")

```



✓ Saved: seasonal_differencing.png

8. ACF and PACF Analysis

Mathematical Background

Autocorrelation Function (ACF):

$$\rho(k) = \frac{\text{Cov}(y * t, y * t - k)}{\text{Var}(y_t)}$$

Partial Autocorrelation Function (PACF):

The correlation between $y * t$ and $y * t - k$ after removing the effect of intermediate lags.

Model Identification:

- **AR(p)**: PACF cuts off after lag p , ACF decays
- **MA(q)**: ACF cuts off after lag q , PACF decays
- **ARMA(p,q)**: Both ACF and PACF decay

```
In [15]: # ACF and PACF plots
fig, axes = plt.subplots(2, 2, figsize=(16, 10))

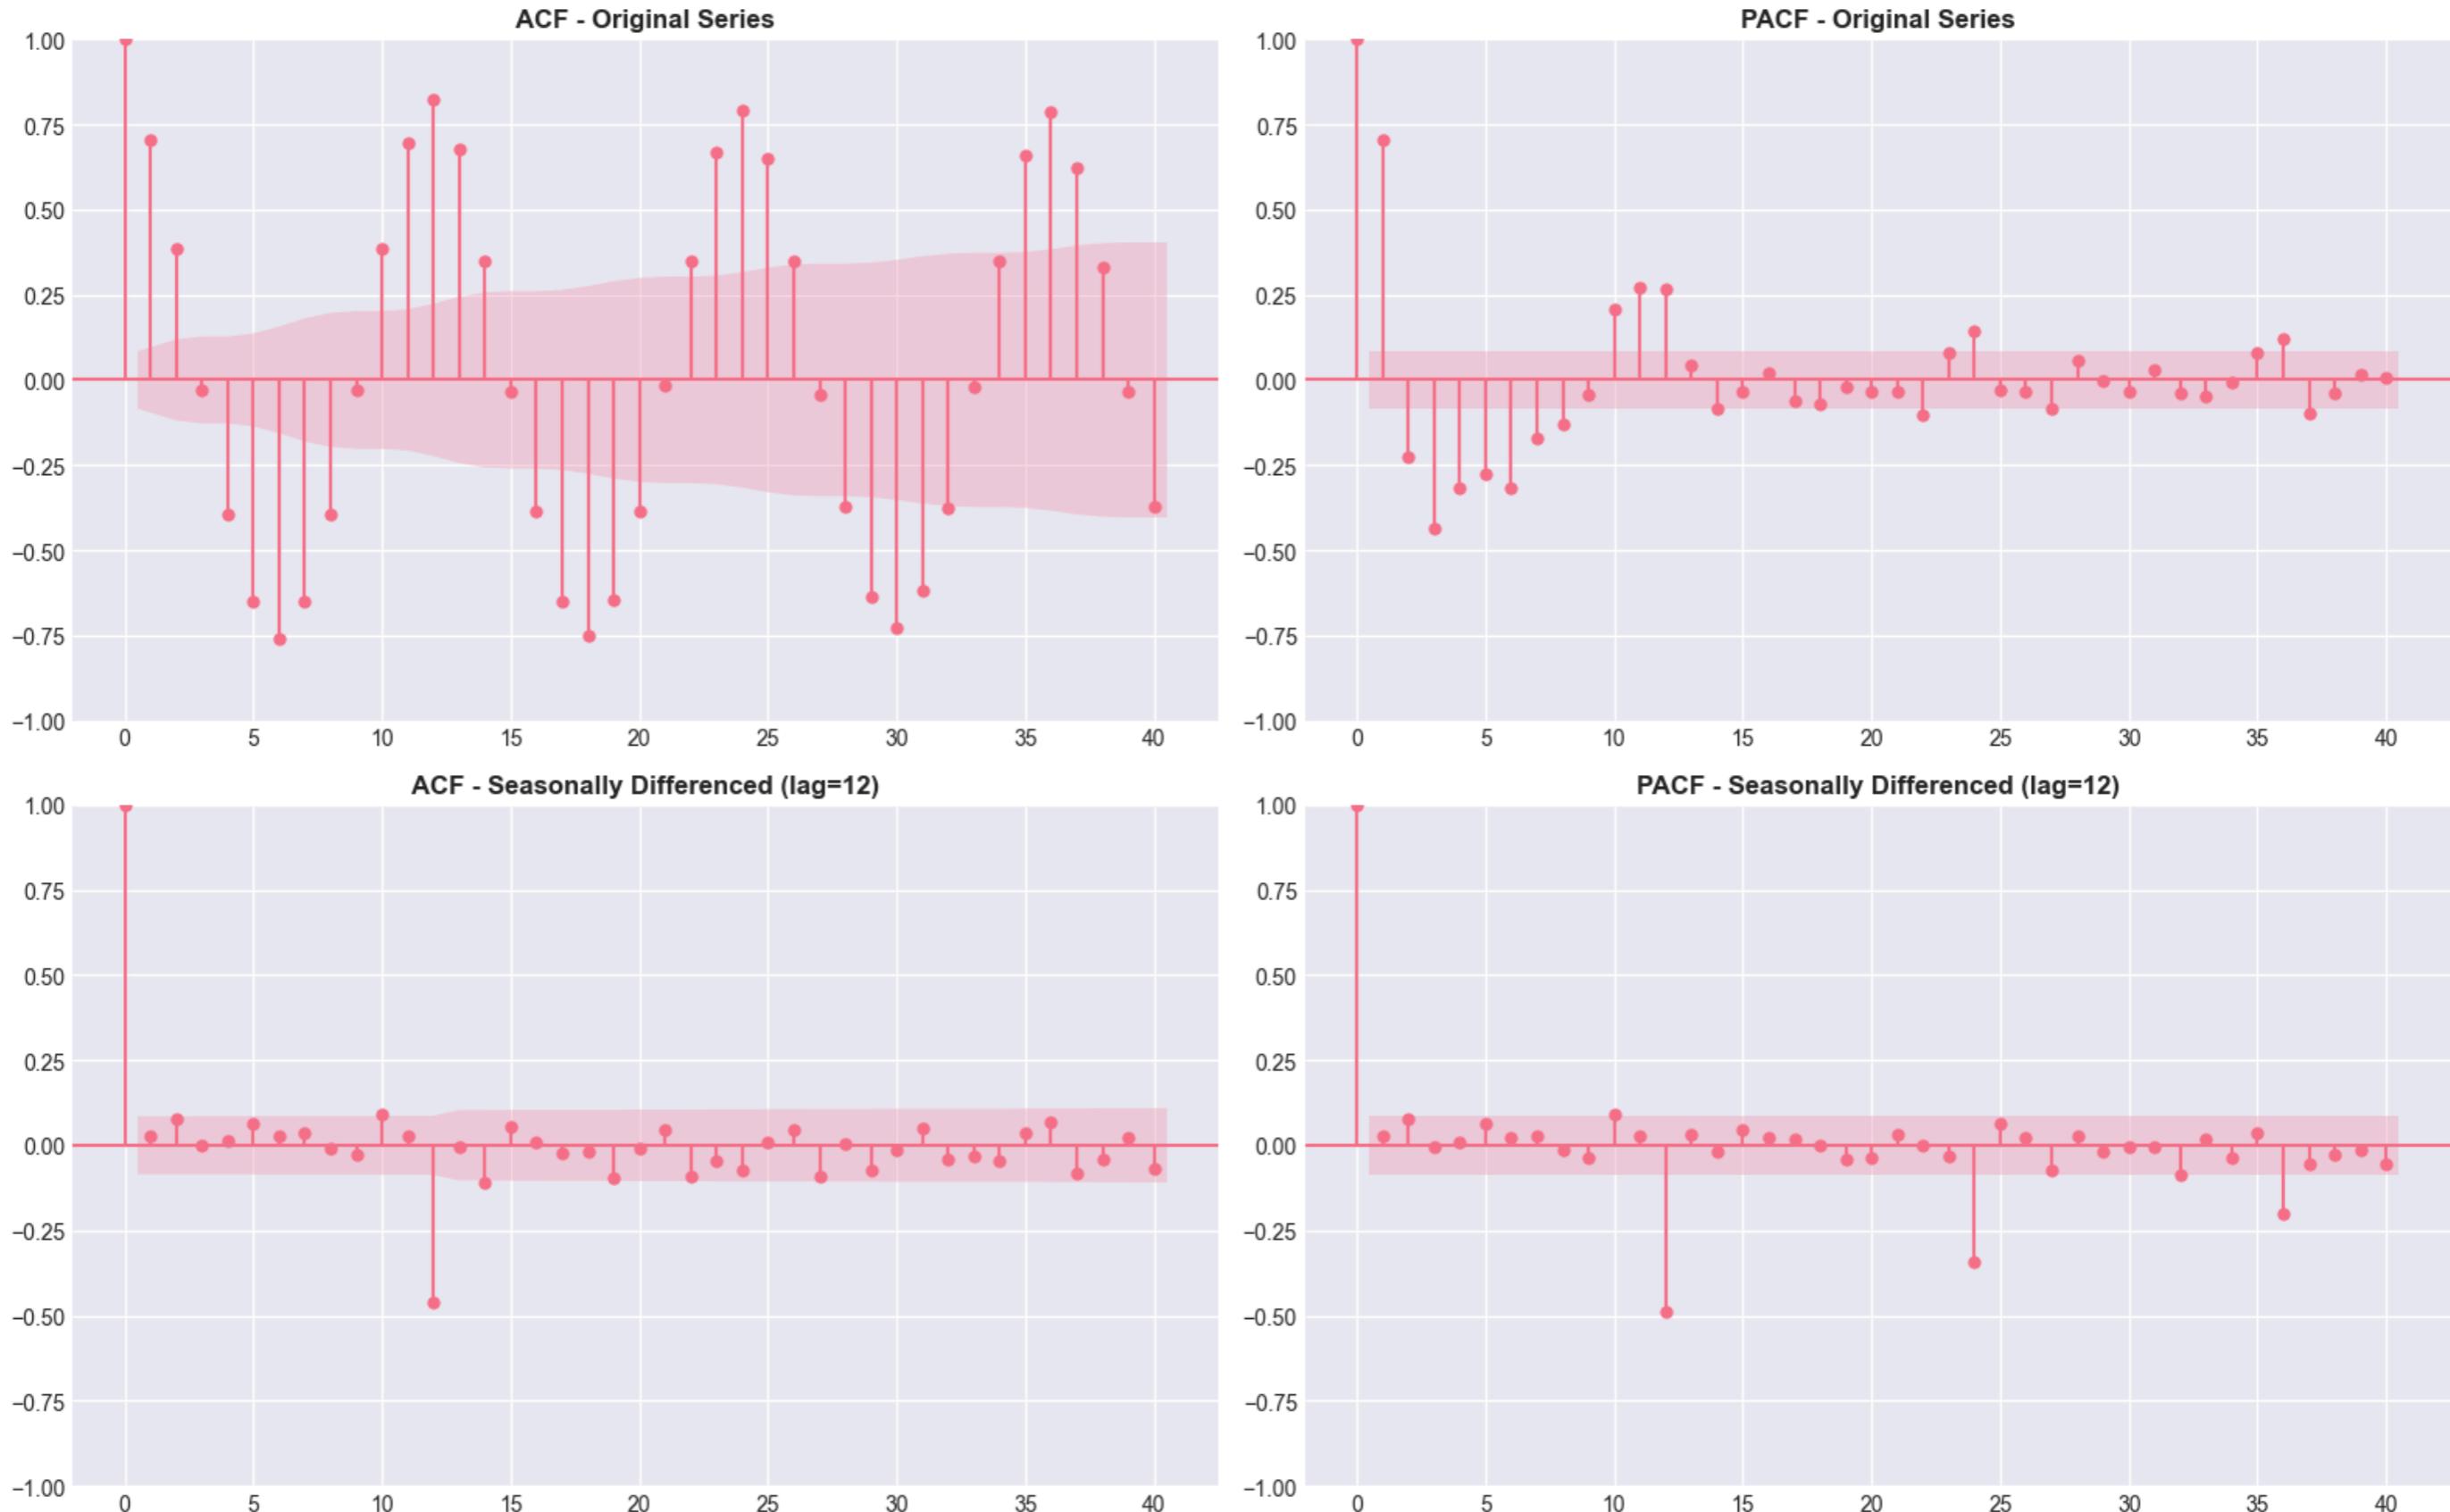
# Original series
plot_acf(ts_data.dropna(), lags=40, ax=axes[0, 0])
axes[0, 0].set_title("ACF - Original Series", fontsize=12, fontweight="bold")

plot_pacf(ts_data.dropna(), lags=40, ax=axes[0, 1])
axes[0, 1].set_title("PACF - Original Series", fontsize=12, fontweight="bold")

# Seasonally differenced series
plot_acf(ts_seasonal_diff.dropna(), lags=40, ax=axes[1, 0])
axes[1, 0].set_title(
    "ACF - Seasonally Differenced (lag=12)", fontsize=12, fontweight="bold"
)

plot_pacf(ts_seasonal_diff.dropna(), lags=40, ax=axes[1, 1])
axes[1, 1].set_title(
    "PACF - Seasonally Differenced (lag=12)", fontsize=12, fontweight="bold"
)

plt.tight_layout()
plt.savefig("acf_pacf.png", dpi=300, bbox_inches="tight")
plt.show()
print("✓ Saved: acf_pacf.png")
```



✓ Saved: acf_pacf.png

9. Train-Test Split (No Data Leakage)

Correct Approach: Temporal Split

To prevent **data leakage**, we split the data temporally:

- **Train set:** Earlier observations (for model fitting)
- **Test set:** Later observations (for out-of-sample evaluation)

We use the **last 60 months (~5 years)** as the test set.

```
In [16]: # Define test horizon
test_horizon = 60 # Last 5 years for testing

# Split data
train_data = ts_data.iloc[: -test_horizon]
test_data = ts_data.iloc[-test_horizon :]

print("=" * 60)
print("Train-Test Split")
print("=" * 60)
print(
    f"Train: {train_data.index[0]} to {train_data.index[-1]} ({len(train_data)} months)"
)
print(f"Test: {test_data.index[0]} to {test_data.index[-1]} ({len(test_data)} months)")
print(f"\nTrain set: {len(train_data)} observations")
print(f"Test set: {len(test_data)} observations")
print(f"Total: {len(ts_data)} observations")
```

```
=====
Train-Test Split
=====
Train: 1981-01-01 00:00:00 to 2021-01-01 00:00:00 (481 months)
Test: 2021-02-01 00:00:00 to 2026-01-01 00:00:00 (60 months)
```

```
Train set: 481 observations
Test set: 60 observations
Total: 541 observations
```

```
In [17]: # Visualize train-test split
fig, ax = plt.subplots(figsize=(15, 6))

ax.plot(
    train_data.index,
    train_data.values,
    linewidth=1.5,
    color="blue",
    label="Train",
    alpha=0.8,
)
ax.plot(
    test_data.index,
    test_data.values,
    linewidth=1.5,
    color="red",
    label="Test",
    alpha=0.8,
)
ax.axvline(x=test_data.index[0], color="black", linestyle="--", linewidth=2, alpha=0.5)

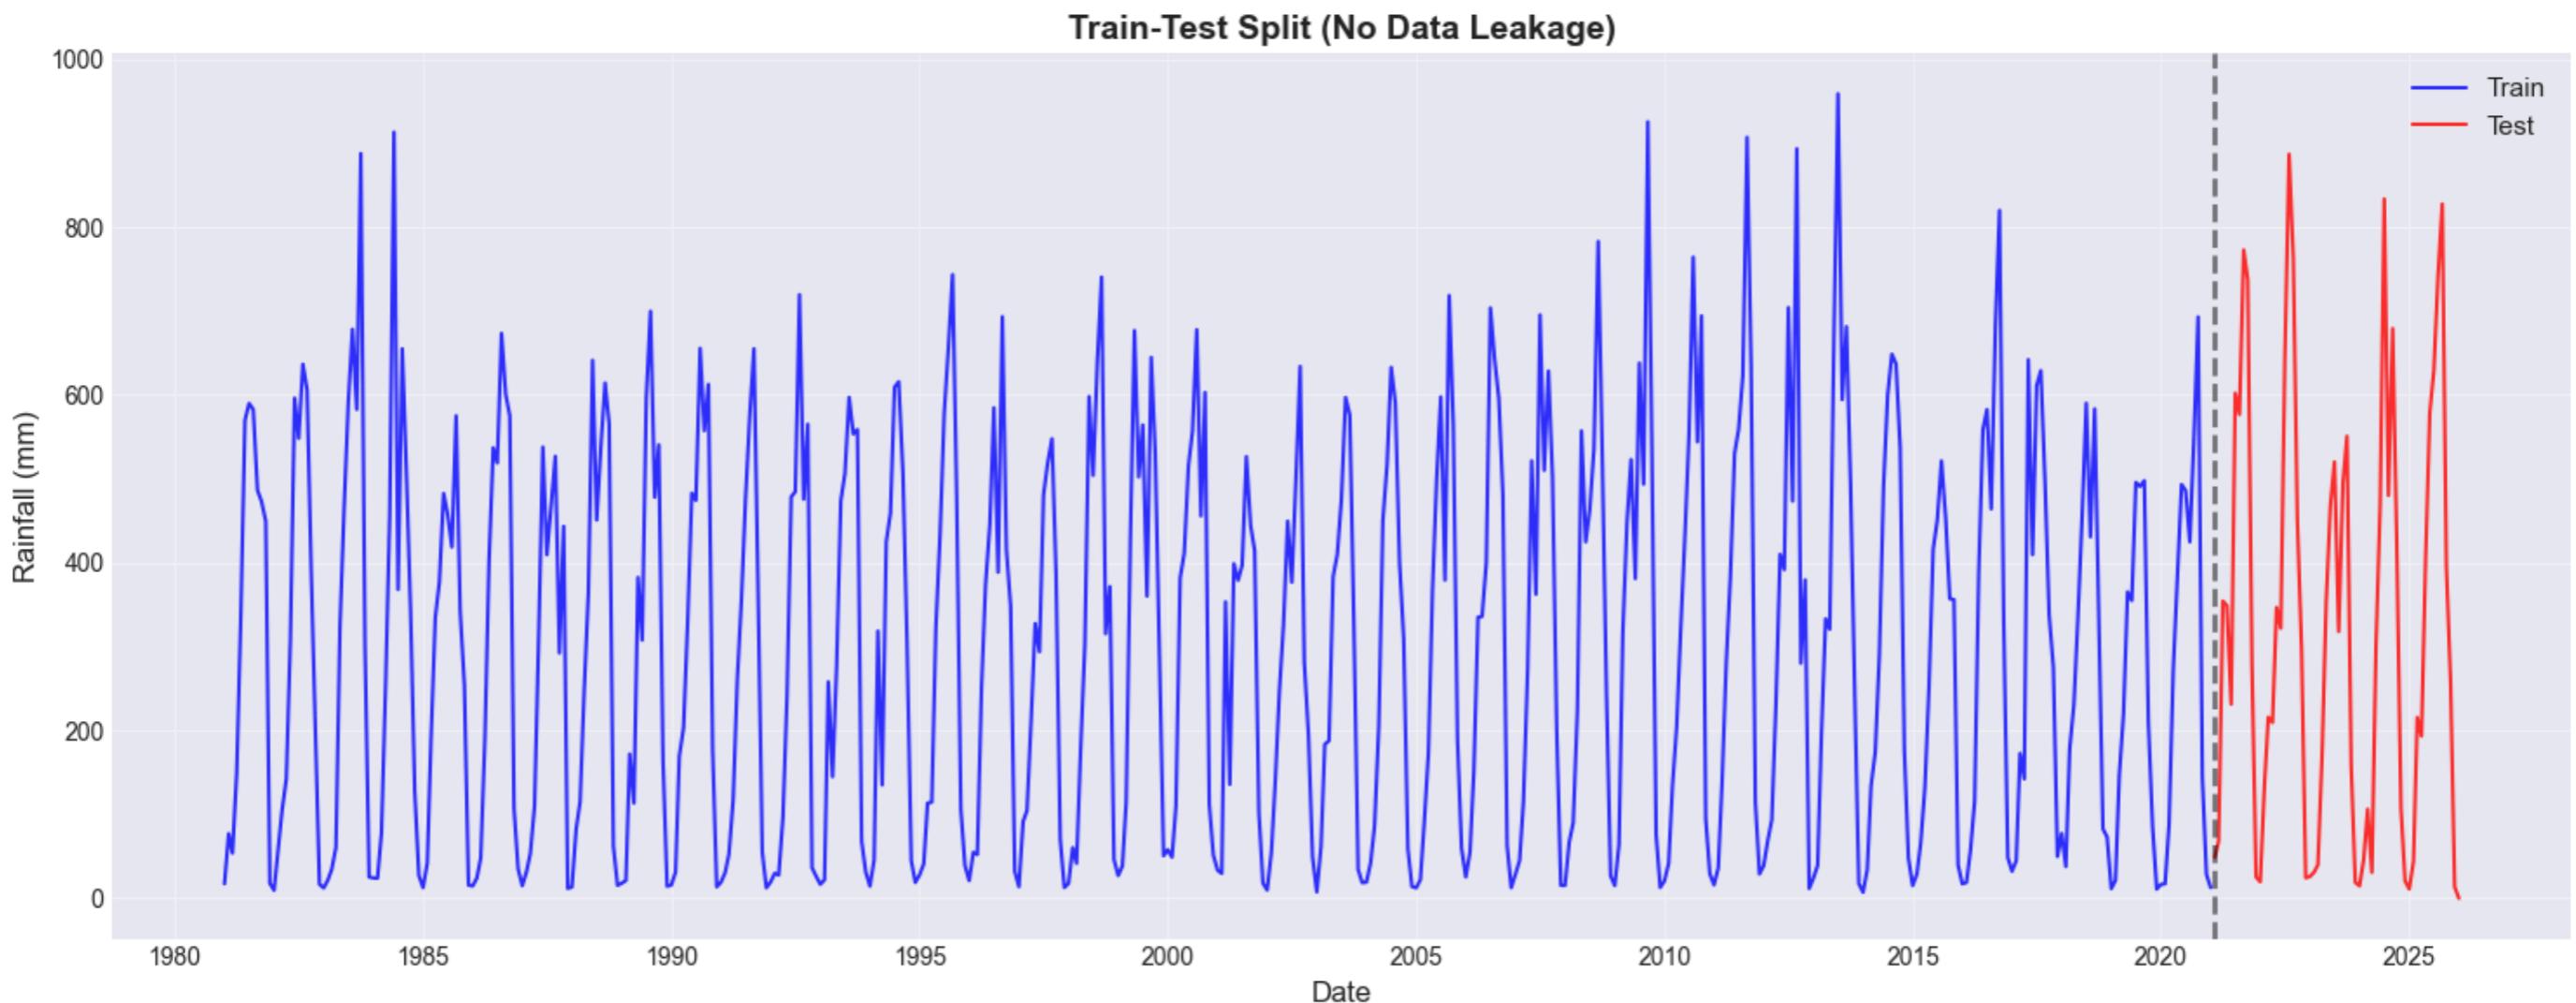
ax.set_title("Train-Test Split (No Data Leakage)", fontsize=14, fontweight="bold")
ax.set_xlabel("Date", fontsize=12)
ax.set_ylabel("Rainfall (mm)", fontsize=12)
```

```

ax.legend(fontsize=11)
ax.grid(True, alpha=0.3)

plt.tight_layout()
plt.savefig("train_test_split.png", dpi=300, bbox_inches="tight")
plt.show()
print("✓ Saved: train_test_split.png")

```



✓ Saved: train_test_split.png

10. Model Fitting with Out-of-Sample Evaluation

Correct Approach: Rolling Forecast

We use **rolling one-step-ahead forecasts** on the test set:

1. Fit model on train data
2. Forecast 1 step ahead
3. Add actual observation to training set
4. Repeat for entire test period

This provides **true out-of-sample evaluation** without data leakage.

Modern API: No `disp` parameter

All models use:

- `ARIMA()` from `statsmodels.tsa.arima.model`
- `SARIMAX()` from `statsmodels.tsa.statespace.sarimax`
- `.fit()` with **NO** `disp` parameter (removed in statsmodels 0.14+)

```
In [18]: def rolling_forecast_evaluation(series, order, seasonal_order=None, model_name="Model"):  
    """  
    Perform rolling one-step-ahead forecast on test set.  
  
    Parameters:  
    -----  
    series : pd.Series  
        Full time series (train + test)  
    order : tuple  
        ARIMA order (p, d, q)  
    seasonal_order : tuple or None  
        SARIMA seasonal order (P, D, Q, s) or None  
    model_name : str  
        Name of the model for reporting  
  
    Returns:  
    -----  
    predictions : pd.Series  
        Out-of-sample predictions  
    mae : float  
        Mean Absolute Error  
    rmse : float  
        Root Mean Squared Error  
    """  
  
    # Split indices  
    train_idx = series.index[: -test_horizon]  
    test_idx = series.index[-test_horizon:]  
  
    y_train = series.loc[train_idx]  
    y_test = series.loc[test_idx]  
  
    # Rolling forecast  
    history = y_train.copy()  
    predictions = []  
  
    print(f"\nFitting {model_name}...")  
    print(f" Order: {order}")  
    if seasonal_order:  
        print(f" Seasonal Order: {seasonal_order}")  
  
    for t in range(len(y_test)):  
        try:  
            # Fit model - MODERN API (no disp parameter)  
            if seasonal_order is None:  
                model = ARIMA(history, order=order)  
            else:  
                model = SARIMAX(  
                    history,  
                    order=order,  
                    seasonal_order=seasonal_order,
```

```

        enforce_stationarity=False,
        enforce_invertibility=False,
    )

    # Fit without deprecated 'disp' parameter
    fitted = model.fit()

    # One-step forecast
    forecast = fitted.forecast(steps=1)
    pred_value = max(0, forecast.iloc[0]) # Ensure non-negative
    predictions.append(pred_value)

    # Add actual observation to history
    history = pd.concat(
        [history, pd.Series([y_test.iloc[t]], index=[y_test.index[t]])]
    )

except Exception as e:
    print(f" X Error at step {t}: {str(e)[:50]}")
    predictions.append(history.mean()) # Fallback to mean

# Convert to series
predictions = pd.Series(predictions, index=test_idx)

# Calculate metrics
mae = mean_absolute_error(y_test, predictions)
rmse = np.sqrt(mean_squared_error(y_test, predictions))

print(f" ✓ MAE: {mae:.2f}")
print(f" ✓ RMSE: {rmse:.2f}")

return predictions, mae, rmse

# Store results
model_results = []

```

10.1 AR Models

Autoregressive (AR) models:

$$y * t = c + \phi_1 y * t - 1 + \phi_2 y * t - 2 + \dots + \phi_p y * t - p + \varepsilon_t$$

Equivalent to ARIMA($p, 0, 0$)

```
In [19]: print("=" * 60)
print("AR MODELS (Autoregressive)")
print("=" * 60)

ar_orders = [
    (1, 0, 0), # AR(1)
    (2, 0, 0), # AR(2)
    (3, 0, 0), # AR(3)
]

for p, d, q in ar_orders:
    name = f"AR({p})"
    preds, mae, rmse = rolling_forecast_evaluation(
```

```

        ts_data, order=(p, d, q), seasonal_order=None, model_name=name
    )
    model_results.append(
        {
            "Model": name,
            "Order": (p, d, q),
            "Seasonal Order": None,
            "MAE": mae,
            "RMSE": rmse,
            "Predictions": preds,
        }
    )
print("\n✓ AR models completed")
=====
```

AR MODELS (Autoregressive)

Fitting AR(1)...

Order: (1, 0, 0)
✓ MAE: 150.95
✓ RMSE: 180.45

Fitting AR(2)...

Order: (2, 0, 0)
✓ MAE: 138.95
✓ RMSE: 171.25

Fitting AR(3)...

Order: (3, 0, 0)
✓ MAE: 125.66
✓ RMSE: 160.05

✓ AR models completed

10.2 MA Models

Moving Average (MA) models:

$$y_t = \mu + \varepsilon_t + \theta_1 \varepsilon_{t-1} + \theta_2 \varepsilon_{t-2} + \dots + \theta_q \varepsilon_{t-q}$$

Equivalent to ARIMA(0, 0, q)

```
In [20]: print("=" * 60)
print("MA MODELS (Moving Average)")
print("=" * 60)

ma_orders = [
    (0, 0, 1), # MA(1)
    (0, 0, 2), # MA(2)
    (0, 0, 3), # MA(3)
]

for p, d, q in ma_orders:
    name = f"MA({q})"
    preds, mae, rmse = rolling_forecast_evaluation(
        ts_data, order=(p, d, q), seasonal_order=None, model_name=name
    )
```

```

model_results.append(
{
    "Model": name,
    "Order": (p, d, q),
    "Seasonal Order": None,
    "MAE": mae,
    "RMSE": rmse,
    "Predictions": preds,
}
)
print("\n✓ MA models completed")
=====
```

MA MODELS (Moving Average)

Fitting MA(1)...

Order: (0, 0, 1)
✓ MAE: 164.69
✓ RMSE: 197.02

Fitting MA(2)...

Order: (0, 0, 2)
✓ MAE: 150.98
✓ RMSE: 184.62

Fitting MA(3)...

Order: (0, 0, 3)
✓ MAE: 141.34
✓ RMSE: 171.19

✓ MA models completed

10.3 ARMA Models

ARMA (Autoregressive Moving Average) models:

$$y * t = c + \phi_1 y * t - 1 + \dots + \phi * p y * t - p + \varepsilon * t + \theta_1 \varepsilon * t - 1 + \dots + \theta * q \varepsilon * t - q$$

Equivalent to ARIMA($p, 0, q$)

```
In [35]: print("=" * 60)
print("ARMA MODELS (Autoregressive Moving Average)")
print("=" * 60)

arma_orders = [
    (1, 0, 1), # ARMA(1,1)
    (2, 0, 1), # ARMA(2,1)
    (1, 0, 2), # ARMA(1,2)
    (2, 0, 2), # ARMA(2,2)
]

for p, d, q in arma_orders:
    name = f"ARMA({p},{q})"
    preds, mae, rmse = rolling_forecast_evaluation(
        ts_data, order=(p, d, q), seasonal_order=None, model_name=name
    )
    model_results.append(
```

```

    }
    "Model": name,
    "Order": (p, d, q),
    "Seasonal Order": None,
    "MAE": mae,
    "RMSE": rmse,
    "Predictions": preds,
)
)

print("\n✓ ARMA models completed")
=====

ARMA MODELS (Autoregressive Moving Average)
=====

Fitting ARMA(1,1)...
Order: (1, 0, 1)
✓ MAE: 145.67
✓ RMSE: 175.96

Fitting ARMA(2,1)...
Order: (2, 0, 1)
✓ MAE: 112.53
✓ RMSE: 144.33

Fitting ARMA(1,2)...
Order: (1, 0, 2)
✓ MAE: 140.35
✓ RMSE: 172.53

Fitting ARMA(2,2)...
Order: (2, 0, 2)
✓ MAE: 89.78
✓ RMSE: 126.30

✓ ARMA models completed

```

10.4 ARIMA Models

ARIMA (Autoregressive Integrated Moving Average) models:

$$(1 - \phi_1 B - \cdots - \phi_p B^p)(1 - B)^d y_t = (1 + \theta_1 B + \cdots + \theta_q B^q) \varepsilon_t$$

where B is the backshift operator and d is the differencing order.

```
In [36]: print("=" * 60)
print("ARIMA MODELS (Autoregressive Integrated Moving Average)")
print("=" * 60)

arima_orders = [
    (1, 0, 1), # ARIMA(1,0,1)
    (2, 0, 1), # ARIMA(2,0,1)
    (1, 1, 1), # ARIMA(1,1,1)
    (2, 1, 1), # ARIMA(2,1,1)
    (1, 0, 2), # ARIMA(1,0,2)
]

for p, d, q in arima_orders:
```

```

name = f"ARIMA({p},{d},{q})"
preds, mae, rmse = rolling_forecast_evaluation(
    ts_data, order=(p, d, q), seasonal_order=None, model_name=name
)
model_results.append(
{
    "Model": name,
    "Order": (p, d, q),
    "Seasonal Order": None,
    "MAE": mae,
    "RMSE": rmse,
    "Predictions": preds,
}
)
print("\n✓ ARIMA models completed")

```

=====
ARIMA MODELS (Autoregressive Integrated Moving Average)
=====

Fitting ARIMA(1,0,1)...

Order: (1, 0, 1)
✓ MAE: 145.67
✓ RMSE: 175.96

Fitting ARIMA(2,0,1)...

Order: (2, 0, 1)
✓ MAE: 112.53
✓ RMSE: 144.33

Fitting ARIMA(1,1,1)...

Order: (1, 1, 1)
✓ MAE: 150.18
✓ RMSE: 195.07

Fitting ARIMA(2,1,1)...

Order: (2, 1, 1)
✓ MAE: 145.03
✓ RMSE: 192.92

Fitting ARIMA(1,0,2)...

Order: (1, 0, 2)
✓ MAE: 140.35
✓ RMSE: 172.53

✓ ARIMA models completed

10.5 SARIMA Models

SARIMA (Seasonal ARIMA) models:

$$\Phi_P(B^s)\phi_p(B)\nabla_s^D\nabla^d y_t = \Theta_Q(B^s)\theta_q(B)\varepsilon_t$$

where:

- (p, d, q) = Non-seasonal order
- (P, D, Q, s) = Seasonal order
- $s = 12$ (monthly seasonality)

- $D = 1$ (seasonal differencing applied)

```
In [23]: print("=" * 60)
print("SARIMA MODELS (Seasonal ARIMA)")
print("=" * 60)

sarima_specs = [
    ((1, 0, 1), (1, 1, 1, 12)), # SARIMA(1,0,1)(1,1,1)[12]
    ((2, 0, 1), (1, 1, 1, 12)), # SARIMA(2,0,1)(1,1,1)[12]
    ((1, 1, 1), (1, 1, 1, 12)), # SARIMA(1,1,1)(1,1,1)[12]
    ((0, 0, 1), (1, 1, 1, 12)), # SARIMA(0,0,1)(1,1,1)[12]
    ((1, 0, 0), (1, 1, 0, 12)), # SARIMA(1,0,0)(1,1,0)[12]
]

for order, seasonal in sarima_specs:
    p, d, q = order
    P, D, Q, s = seasonal
    name = f"SARIMA({p},{d},{q})({P},{D},{Q})[{s}]"

    preds, mae, rmse = rolling_forecast_evaluation(
        ts_data, order=order, seasonal_order=seasonal, model_name=name
    )
    model_results.append(
        {
            "Model": name,
            "Order": order,
            "Seasonal Order": seasonal,
            "MAE": mae,
            "RMSE": rmse,
            "Predictions": preds,
        }
    )

print("\n✓ SARIMA models completed")
```

```
=====
SARIMA MODELS (Seasonal ARIMA)
=====
```

```
Fitting SARIMA(1,0,1)(1,1,1)[12]...
```

```
Order: (1, 0, 1)  
Seasonal Order: (1, 1, 1, 12)  
✓ MAE: 81.01  
✓ RMSE: 110.40
```

```
Fitting SARIMA(2,0,1)(1,1,1)[12]...
```

```
Order: (2, 0, 1)  
Seasonal Order: (1, 1, 1, 12)  
✓ MAE: 80.59  
✓ RMSE: 111.17
```

```
Fitting SARIMA(1,1,1)(1,1,1)[12]...
```

```
Order: (1, 1, 1)  
Seasonal Order: (1, 1, 1, 12)  
✓ MAE: 82.34  
✓ RMSE: 112.15
```

```
Fitting SARIMA(0,0,1)(1,1,1)[12]...
```

```
Order: (0, 0, 1)  
Seasonal Order: (1, 1, 1, 12)  
✓ MAE: 79.63  
✓ RMSE: 109.86
```

```
Fitting SARIMA(1,0,0)(1,1,0)[12]...
```

```
Order: (1, 0, 0)  
Seasonal Order: (1, 1, 0, 12)  
✓ MAE: 105.86  
✓ RMSE: 143.76
```

```
✓ SARIMA models completed
```

11. Model Comparison

Evaluation Metrics

Mean Absolute Error (MAE):

$$\text{MAE} = \frac{1}{n} \sum_{t=1}^n |y_t - \hat{y}_t|$$

Root Mean Squared Error (RMSE):

$$\text{RMSE} = \sqrt{\frac{1}{n} \sum_{t=1}^n (y_t - \hat{y}_t)^2}$$

Lower values indicate better model performance.

```
In [24]: # Create comparison dataframe  
results_df = pd.DataFrame(model_results)  
results_df = results_df[["Model", "Order", "Seasonal Order", "MAE", "RMSE"]].copy()  
results_df = results_df.sort_values("RMSE").reset_index(drop=True)
```

```

results_df["Rank"] = range(1, len(results_df) + 1)

print("=" * 80)
print("MODEL COMPARISON (Ranked by RMSE - Lower is Better)")
print("=" * 80)
print(results_df[["Rank", "Model", "RMSE", "MAE"]].to_string(index=False))

# Save results
results_df.to_csv("model_comparison.csv", index=False)
print("\n✓ Saved: model_comparison.csv")

```

```
=====
MODEL COMPARISON (Ranked by RMSE - Lower is Better)
=====

Rank          Model      RMSE      MAE
1 SARIMA(0,0,1)(1,1,1)[12] 109.858697 79.626403
2 SARIMA(1,0,1)(1,1,1)[12] 110.400450 81.010345
3 SARIMA(2,0,1)(1,1,1)[12] 111.169283 80.593612
4 SARIMA(1,1,1)(1,1,1)[12] 112.153433 82.341398
5 ARMA(2,2)           126.303038 89.782461
6 SARIMA(1,0,0)(1,1,0)[12] 143.758021 105.860428
7 ARMA(2,1)           144.332065 112.531743
8 ARIMA(2,0,1)         144.332065 112.531743
9 AR(3)                160.053640 125.662817
10 MA(3)               171.190868 141.343959
11 AR(2)               171.245785 138.950405
12 ARMA(1,2)           172.532249 140.346265
13 ARIMA(1,0,2)         172.532249 140.346265
14 ARMA(1,1)           175.964035 145.668234
15 ARIMA(1,0,1)         175.964035 145.668234
16 AR(1)                180.451381 150.946721
17 MA(2)               184.624907 150.984915
18 ARIMA(2,1,1)         192.915400 145.029835
19 ARIMA(1,1,1)         195.066665 150.176301
20 MA(1)               197.018972 164.694290
```

✓ Saved: model_comparison.csv

```
In [25]: # Plot model comparison
fig, axes = plt.subplots(1, 2, figsize=(16, 6))

# RMSE comparison
axes[0].barh(results_df["Model"][:10], results_df["RMSE"][:10], color="steelblue")
axes[0].set_xlabel("RMSE", fontsize=12)
axes[0].set_title(
    "Top 10 Models by RMSE (Lower is Better)", fontsize=13, fontweight="bold"
)
axes[0].invert_yaxis()
axes[0].grid(True, alpha=0.3, axis="x")

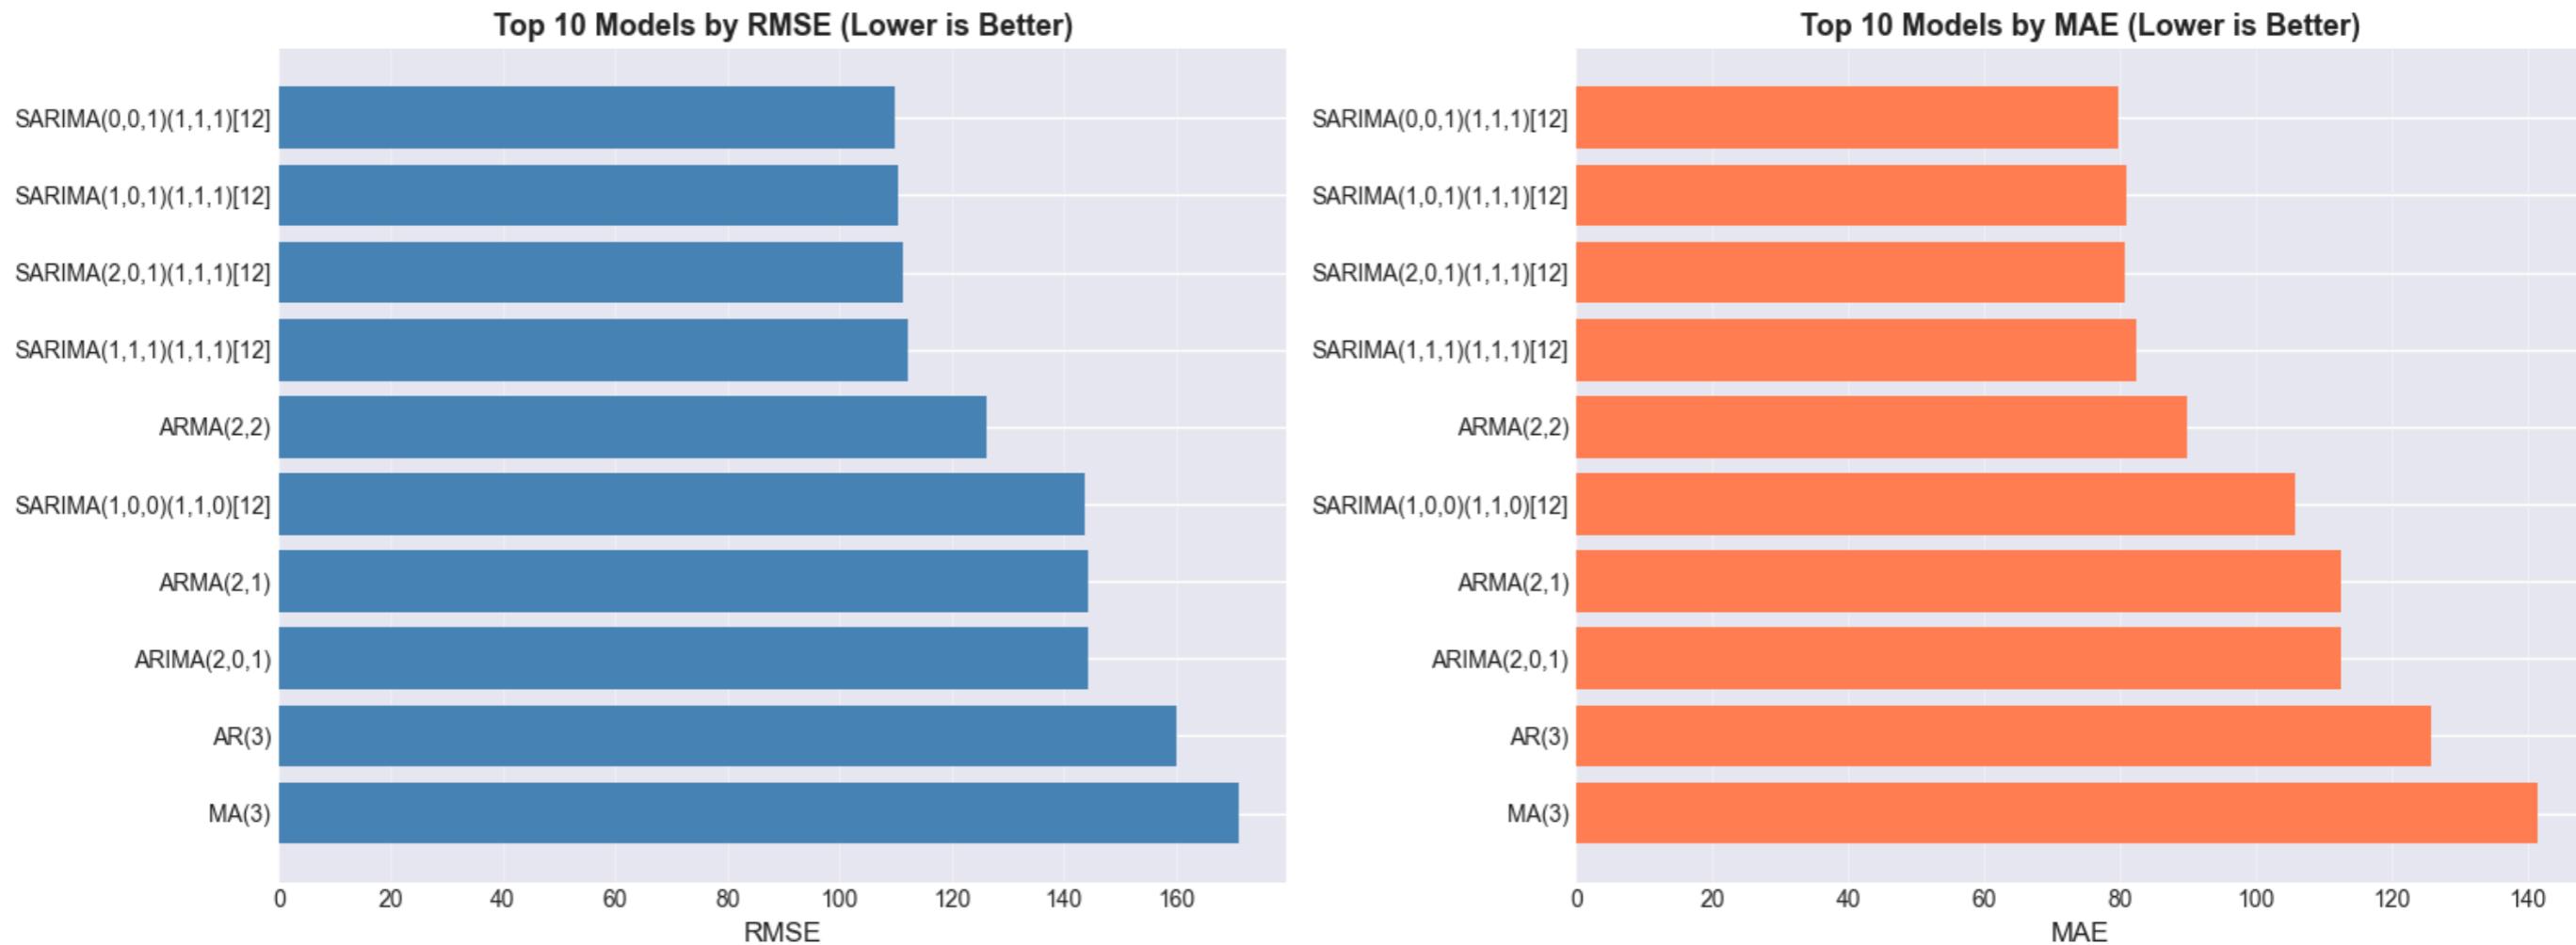
# MAE comparison
axes[1].barh(results_df["Model"][:10], results_df["MAE"][:10], color="coral")
axes[1].set_xlabel("MAE", fontsize=12)
axes[1].set_title(
    "Top 10 Models by MAE (Lower is Better)", fontsize=13, fontweight="bold"
)
axes[1].invert_yaxis()
axes[1].grid(True, alpha=0.3, axis="x")

plt.tight_layout()
```

```

plt.savefig("model_comparison.png", dpi=300, bbox_inches="tight")
plt.show()
print("✓ Saved: model_comparison.png")

```



✓ Saved: model_comparison.png

```

In [26]: # Plot best model forecast
best_model = results_df.iloc[0]
best_name = best_model["Model"]
best_preds = model_results[results_df.index[0]]["Predictions"]

print(f"\n{'='*60}")
print(f"BEST MODEL: {best_name}")
print(f"\n{'='*60}")
print(f"RMSE: {best_model['RMSE']:.2f}")
print(f"MAE: {best_model['MAE']:.2f}")

fig, ax = plt.subplots(figsize=(15, 6))

# Plot last 120 months of train + all test
train_plot = train_data.iloc[-120:]

ax.plot(
    train_plot.index,
    train_plot.values,
    linewidth=1.5,
    color="blue",
    label="Train",
    alpha=0.7,
)

```

```

ax.plot(
    test_data.index,
    test_data.values,
    linewidth=2,
    color="black",
    label="Actual Test",
    marker="o",
    markersize=3,
)
ax.plot(
    best_preds.index,
    best_preds.values,
    linewidth=2,
    color="red",
    label=f"{best_name} Forecast",
    linestyle="--",
    marker="s",
    markersize=3,
)

ax.axvline(x=test_data.index[0], color="gray", linestyle="--", linewidth=2, alpha=0.5)
ax.set_title(
    f'Best Model: {best_name} (RMSE={best_model["RMSE"]:.2f})',
    fontsize=14,
    fontweight="bold",
)
ax.set_xlabel("Date", fontsize=12)
ax.set_ylabel("Rainfall (mm)", fontsize=12)
ax.legend(fontsize=11)
ax.grid(True, alpha=0.3)

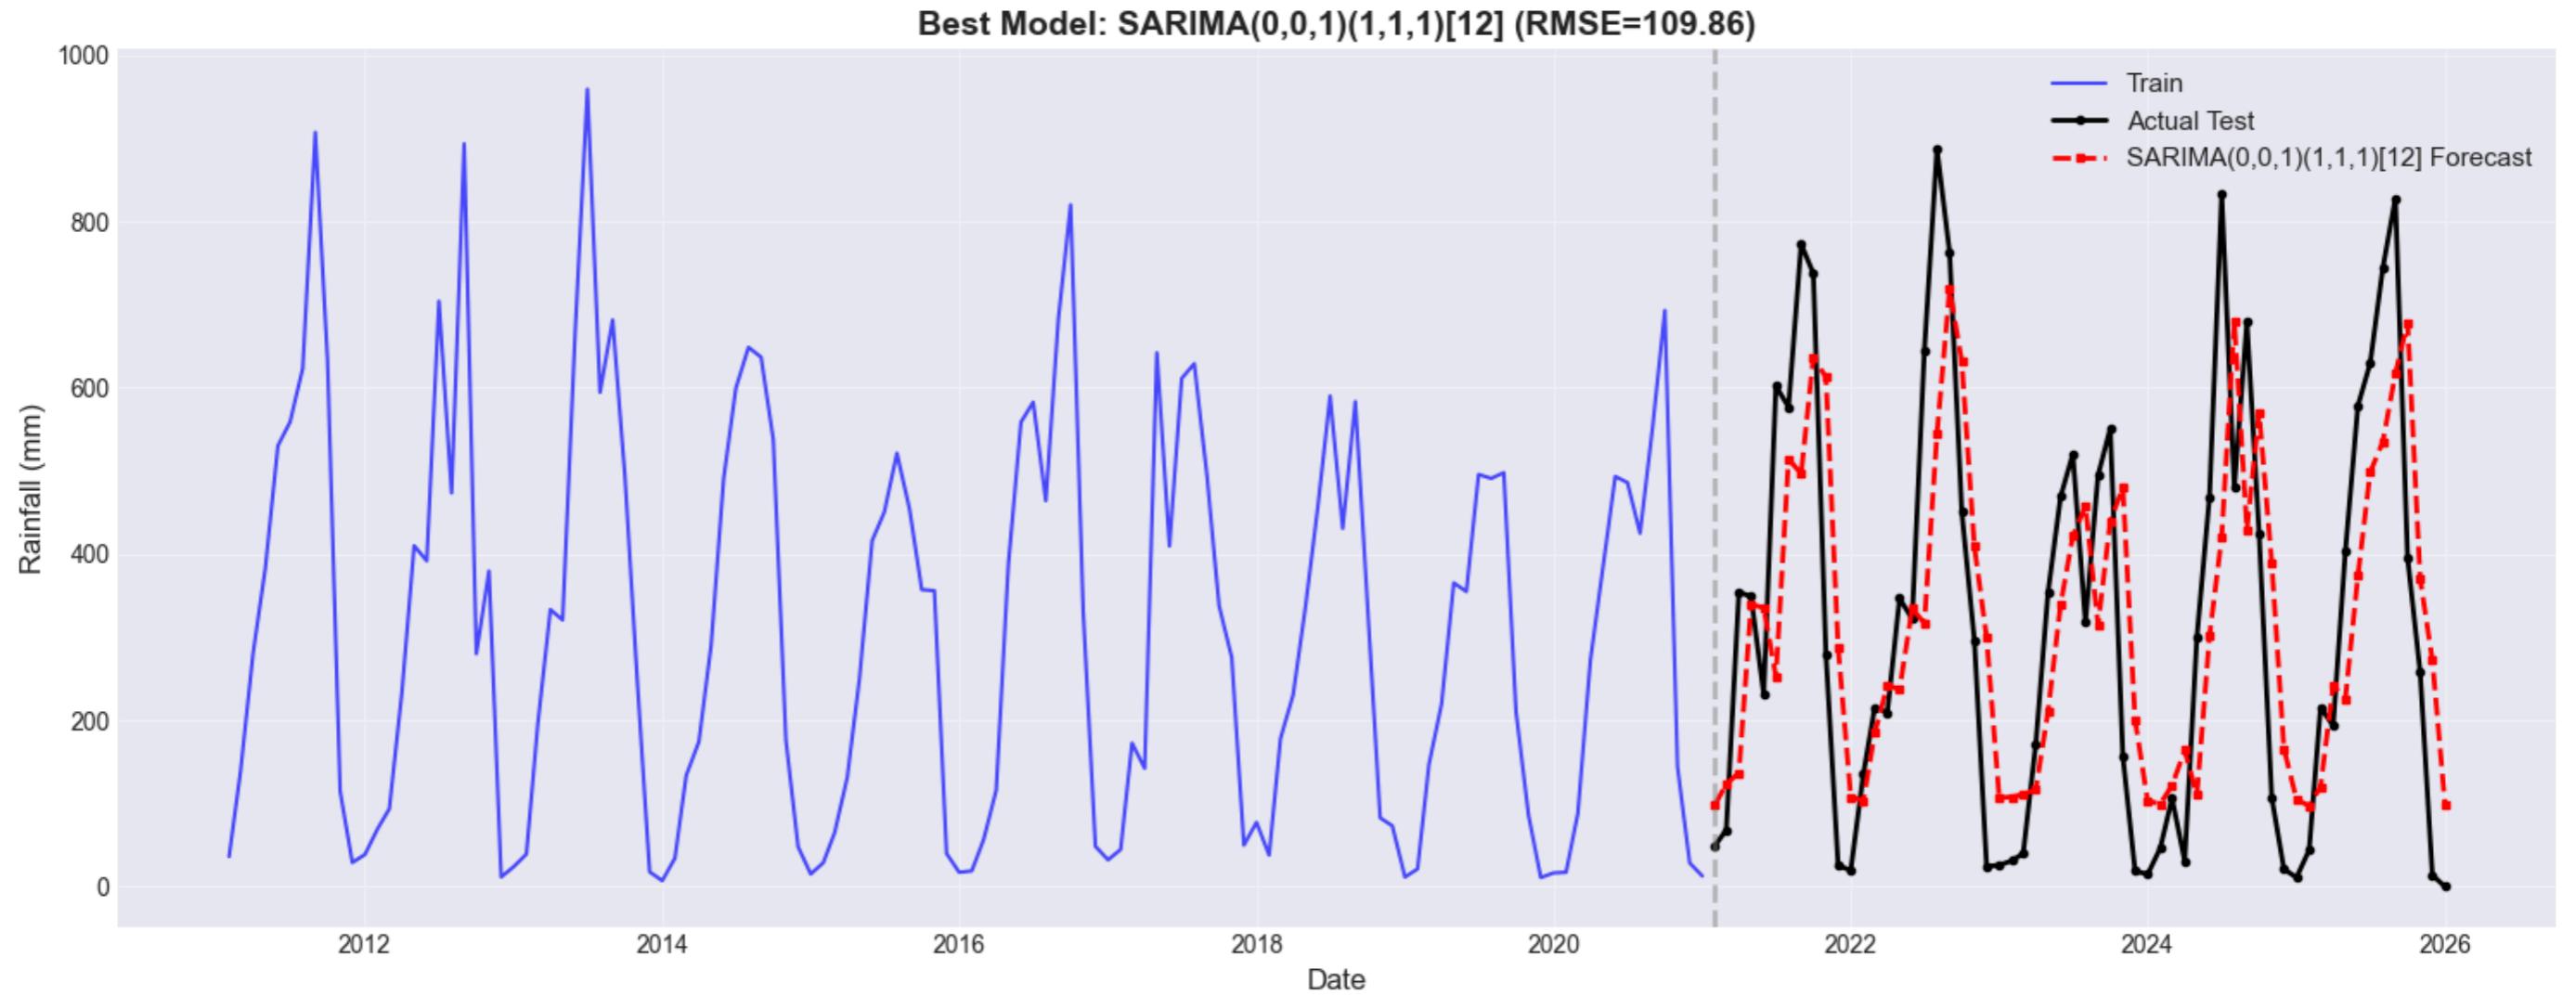
plt.tight_layout()
plt.savefig("best_model_forecast.png", dpi=300, bbox_inches="tight")
plt.show()
print("✓ Saved: best_model_forecast.png")

```

```

=====
BEST MODEL: SARIMA(0,0,1)(1,1,1)[12]
=====
RMSE: 109.86
MAE: 79.63

```



✓ Saved: best_model_forecast.png

12. Residual Diagnostics

Mathematical Background: Residual Analysis

Residuals: $e_t = y_t - \hat{y}_t$

Good model should have:

1. **Zero mean:** $\mathbb{E}[e_t] = 0$
2. **Constant variance:** $\text{Var}(e_t) = \sigma^2$ (homoscedasticity)
3. **No autocorrelation:** $\text{Cov}(e_t, e_s) = 0$ for $t \neq s$
4. **Normality:** $e_t \sim \mathcal{N}(0, \sigma^2)$

Ljung-Box Test:

- H_0 : No autocorrelation in residuals
- H_1 : Autocorrelation present
- Decision: $p\text{-value} > 0.05 \rightarrow$ Good model

```
In [27]: # Fit best model on full train data for diagnostics
best_order = best_model["Order"]
best_seasonal = best_model["Seasonal Order"]

print(f"\n{'='*60}")
print("RESIDUAL DIAGNOSTICS: {best_name}")
print(f"\n{'='*60}")

# Fixed conditional check
if best_seasonal is None or not isinstance(best_seasonal, tuple):
    # Non-seasonal ARIMA model
    final_model = ARIMA(train_data, order=best_order)
else:
    # Seasonal SARIMA model
    final_model = SARIMAX(
        train_data,
        order=best_order,
        seasonal_order=best_seasonal,
        enforce_stationarity=False,
        enforce_invertibility=False,
    )

final_fitted = final_model.fit()
print(final_fitted.summary())
```

```
=====
RESIDUAL DIAGNOSTICS: SARIMA(0,0,1)(1,1,1)[12]
=====

SARIMAX Results
=====

Dep. Variable: rainfall No. Observations: 481
Model: SARIMAX(0, 0, 1)x(1, 1, 1, 12) Log Likelihood -2743.282
Date: Fri, 09 Jan 2026 AIC 5494.565
Time: 20:18:35 BIC 5511.046
Sample: 01-01-1981 HQIC 5501.058
- 01-01-2021
Covariance Type: opg
=====

            coef  std err      z   P>|z|   [0.025   0.975]
-----
ma.L1      0.0399  0.040    1.006    0.315   -0.038   0.118
ar.S.L12   0.0183  0.041    0.448    0.654   -0.062   0.098
ma.S.L12  -1.0897  0.028   -39.191   0.000   -1.144  -1.035
sigma2     8235.6555 580.311   14.192   0.000   7098.267 9373.044
=====

Ljung-Box (L1) (Q): 0.01 Jarque-Bera (JB): 102.16
Prob(Q): 0.91 Prob(JB): 0.00
Heteroskedasticity (H): 1.28 Skew: 0.66
Prob(H) (two-sided): 0.13 Kurtosis: 4.91
=====

Warnings:
[1] Covariance matrix calculated using the outer product of gradients (complex-step).
```

```
In [39]: # Residual diagnostics plots
residuals = final_fitted.resid

fig, axes = plt.subplots(2, 3, figsize=(18, 10))

# 1. Residual time series
```

```

axes[0, 0].plot(residuals, linewidth=1, color="steelblue")
axes[0, 0].axhline(y=0, color="red", linestyle="--", linewidth=1)
axes[0, 0].set_title("Residuals Over Time", fontsize=12, fontweight="bold")
axes[0, 0].set_xlabel("Date", fontsize=10)
axes[0, 0].set_ylabel("Residuals", fontsize=10)
axes[0, 0].grid(True, alpha=0.3)

# 2. Histogram
axes[0, 1].hist(
    residuals.dropna(), bins=30, edgecolor="black", color="coral", alpha=0.7
)
axes[0, 1].axvline(x=0, color="red", linestyle="--", linewidth=2)
axes[0, 1].set_title("Residual Distribution", fontsize=12, fontweight="bold")
axes[0, 1].set_xlabel("Residuals", fontsize=10)
axes[0, 1].set_ylabel("Frequency", fontsize=10)
axes[0, 1].grid(True, alpha=0.3, axis="y")

# 3. Q-Q plot
from scipy import stats

stats.probplot(residuals.dropna(), dist="norm", plot=axes[0, 2])
axes[0, 2].set_title("Q-Q Plot (Normality Check)", fontsize=12, fontweight="bold")
axes[0, 2].grid(True, alpha=0.3)

# 4. ACF of residuals
plot_acf(residuals.dropna(), lags=40, ax=axes[1, 0])
axes[1, 0].set_title("ACF of Residuals", fontsize=12, fontweight="bold")

# 5. PACF of residuals
plot_pacf(residuals.dropna(), lags=40, ax=axes[1, 1])
axes[1, 1].set_title("PACF of Residuals", fontsize=12, fontweight="bold")

# 6. Ljung-Box test results
lb_test = acorr_ljungbox(residuals.dropna(), lags=[10, 20, 30], return_df=True)
axes[1, 2].axis("off")
axes[1, 2].text(0.1, 0.8, "Ljung-Box Test Results", fontsize=12, fontweight="bold")
axes[1, 2].text(
    0.1, 0.6, f"Lag 10: p-value = {lb_test.loc[10, 'lb_pvalue']:.4f}", fontsize=10
)
axes[1, 2].text(
    0.1, 0.5, f"Lag 20: p-value = {lb_test.loc[20, 'lb_pvalue']:.4f}", fontsize=10
)
axes[1, 2].text(
    0.1, 0.4, f"Lag 30: p-value = {lb_test.loc[30, 'lb_pvalue']:.4f}", fontsize=10
)

if lb_test.loc[20, "lb_pvalue"] > 0.05:
    axes[1, 2].text(
        0.1,
        0.2,
        "✓ No autocorrelation (p > 0.05)",
        fontsize=10,
        color="green",
        fontweight="bold",
    )
else:
    axes[1, 2].text(
        0.1,
        0.2,
        "✗ Autocorrelation detected (p < 0.05)",
        fontsize=10,
        color="red",
        fontweight="bold",
    )

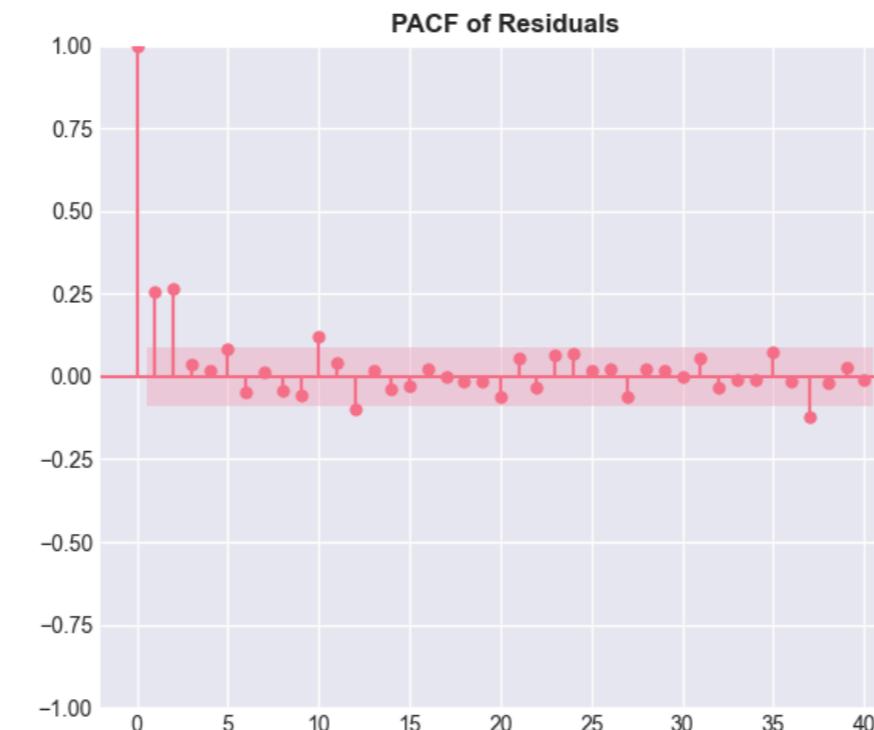
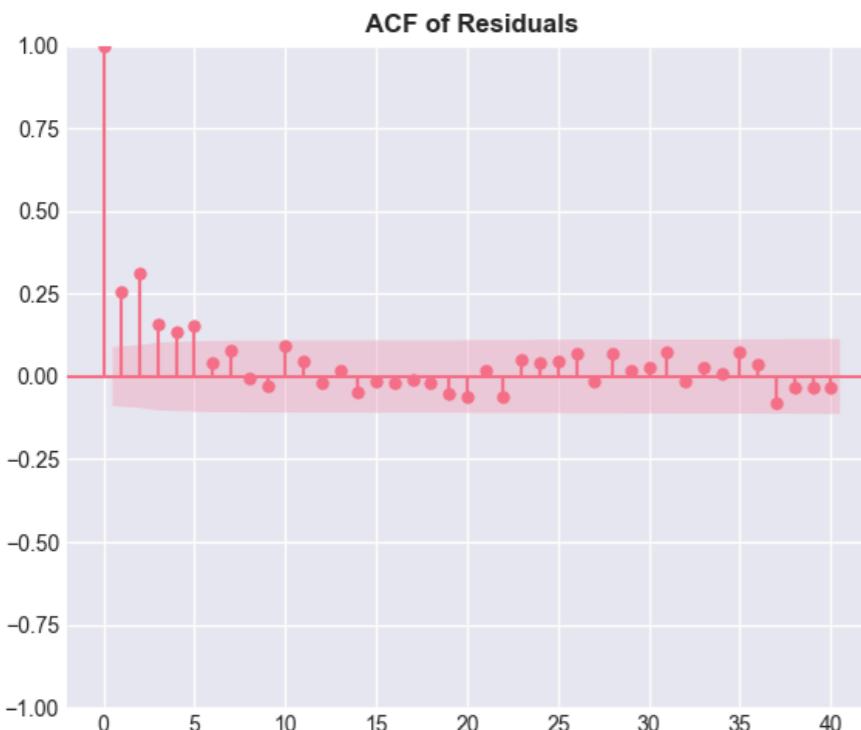
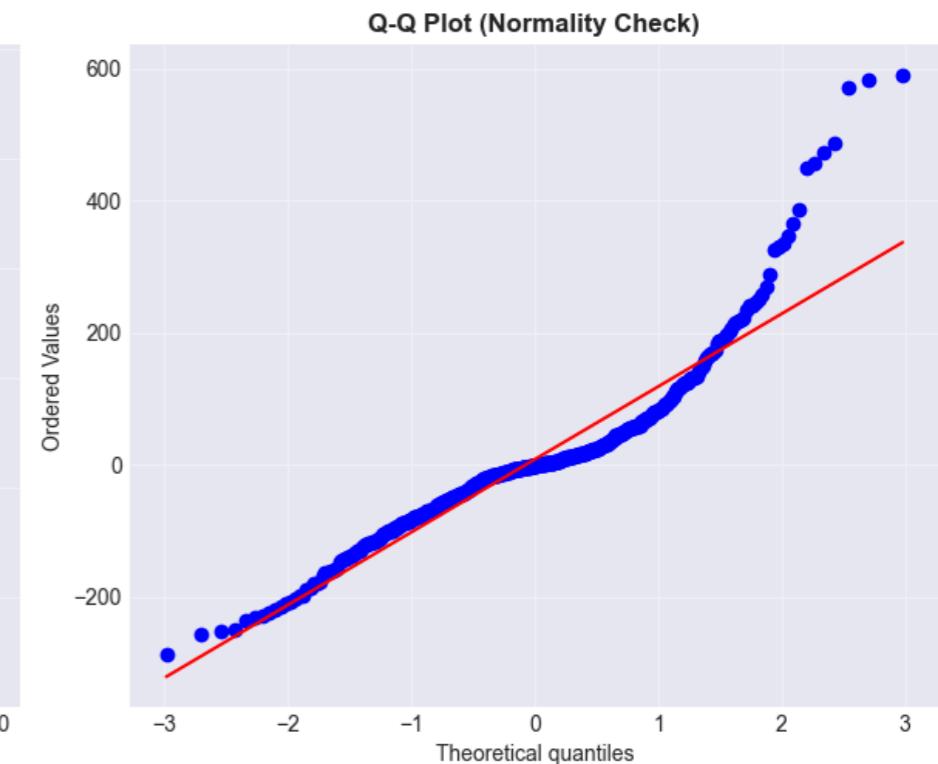
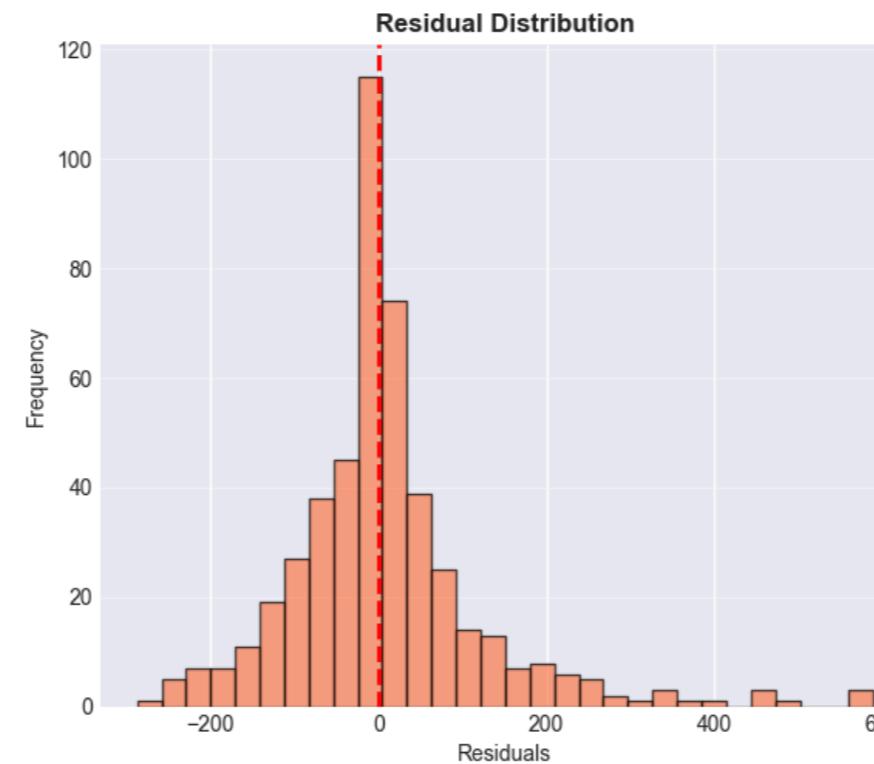
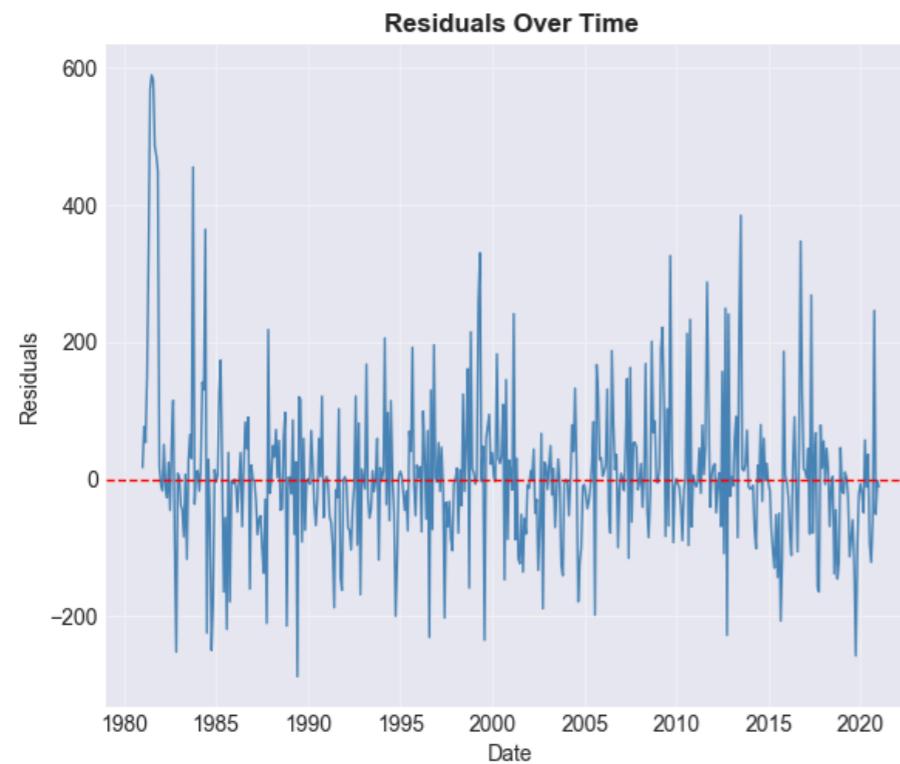
```

```

        fontsize=10,
        color="red",
        fontweight="bold",
    )

plt.tight_layout()
plt.savefig("residual_diagnostics.png", dpi=300, bbox_inches="tight")
plt.show()
print("\n✓ Saved: residual_diagnostics.png")

```



Ljung-Box Test Results

Lag 10: p-value = 0.0000

Lag 20: p-value = 0.0000

Lag 30: p-value = 0.0000

□ Autocorrelation detected (p < 0.05)

✓ Saved: residual_diagnostics.png

13. Final Forecasting

Mathematical Background: h-step Ahead Forecast

$$\hat{y}_{T+h|T} = \mathbb{E}[y_{T+h} | y_1, \dots, y_T]$$

Forecast Variance:

$$\text{Var}(e * T + h|T) = \sigma^2 \sum *i = 0^{h-1} \psi_i^2$$

95% Prediction Interval:

$$\hat{y}_{T+h|T} \pm 1.96 \sqrt{\text{Var}(e_{T+h|T})}$$

```
In [29]: # Refit best model on FULL data for future forecasting
print(f"\n{'='*60}")
print(f"FINAL FORECASTING: {best_name}")
print(f"{'='*60}")

# Fixed conditional check
if best_seasonal is None or not isinstance(best_seasonal, tuple):
    # Non-seasonal ARIMA model
    final_full_model = ARIMA(ts_data, order=best_order)
else:
    # Seasonal SARIMA model
    final_full_model = SARIMAX(
        ts_data,
        order=best_order,
        seasonal_order=best_seasonal,
        enforce_stationarity=False,
        enforce_invertibility=False,
    )

final_full_fitted = final_full_model.fit()

print(f"Model refitted on complete data")
print(f"Observations: {len(ts_data)}")
print(f"\nModel Summary:")
print(final_full_fitted.summary())
```

```
=====
FINAL FORECASTING: SARIMA(0,0,1)(1,1,1)[12]
=====
Model refitted on complete data
Observations: 541

Model Summary:
SARIMAX Results
=====
Dep. Variable: rainfall No. Observations: 541
Model: SARIMAX(0, 0, 1)x(1, 1, 1, 12) Log Likelihood -3110.576
Date: Fri, 09 Jan 2026 AIC 6229.153
Time: 20:18:36 BIC 6246.130
Sample: 01-01-1981 HQIC 6235.806
- 01-01-2026
Covariance Type: opg
=====

            coef  std err      z   P>|z|    [0.025    0.975]
-----
ma.L1      0.0670  0.038    1.777    0.076   -0.007    0.141
ar.S.L12   0.0163  0.037    0.435    0.663   -0.057    0.090
ma.S.L12  -1.0812  0.025   -42.900   0.000   -1.131   -1.032
sigma2     8564.9118 556.152   15.400   0.000  7474.874  9654.950
=====
Ljung-Box (L1) (Q): 0.03 Jarque-Bera (JB): 99.87
Prob(Q): 0.87 Prob(JB): 0.00
Heteroskedasticity (H): 1.24 Skew: 0.63
Prob(H) (two-sided): 0.16 Kurtosis: 4.75
=====
```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
In [30]: # Generate 24-month forecast
forecast_horizon = 24

forecast_result = final_full_fitted.get_forecast(steps=forecast_horizon)
forecast_mean = forecast_result.predicted_mean
forecast_ci = forecast_result.conf_int()

# Ensure non-negative forecasts
forecast_mean = forecast_mean.clip(lower=0)
forecast_ci = forecast_ci.clip(lower=0)

# Create forecast dates
forecast_dates = pd.date_range(
    start=ts_data.index[-1] + pd.DateOffset(months=1),
    periods=forecast_horizon,
    freq="MS",
)

# Create forecast dataframe
forecast_df = pd.DataFrame(
    {
        "date": forecast_dates,
        "forecast": forecast_mean.values,
        "lower_ci": forecast_ci.iloc[:, 0].values,
        "upper_ci": forecast_ci.iloc[:, 1].values,
    }
)
```

```

print(f"\n{'='*60}")
print(f"FUTURE FORECAST ({forecast_horizon} months)")
print(f"\n{'='*60}")
print(forecast_df.to_string(index=False))

forecast_df.to_csv("future_forecast.csv", index=False)
print("\n✓ Saved: future_forecast.csv")

```

```
=====
FUTURE FORECAST (24 months)
=====
```

	date	forecast	lower_ci	upper_ci
2026-02-01	44.489239	0.000000	240.614537	
2026-03-01	134.091103	0.000000	330.653332	
2026-04-01	209.028375	12.466145	405.590604	
2026-05-01	377.382520	180.820291	573.944750	
2026-06-01	453.777624	257.215394	650.339853	
2026-07-01	595.970956	399.408727	792.533186	
2026-08-01	564.600451	368.038222	761.162681	
2026-09-01	648.442053	451.879824	845.004283	
2026-10-01	480.818993	284.256764	677.381223	
2026-11-01	195.159554	0.000000	391.721783	
2026-12-01	27.914789	0.000000	224.477019	
2027-01-01	19.190200	0.000000	215.752424	
2027-02-01	45.707172	0.000000	243.086963	
2027-03-01	132.763310	0.000000	330.144574	
2027-04-01	209.283699	11.902435	406.664964	
2027-05-01	376.965658	179.584394	574.346922	
2027-06-01	451.746374	254.365109	649.127638	
2027-07-01	595.428547	398.047283	792.809811	
2027-08-01	561.686686	364.305422	759.067951	
2027-09-01	645.522288	448.141023	842.903552	
2027-10-01	482.206156	284.824891	679.587420	
2027-11-01	194.139834	0.000000	391.521098	
2027-12-01	28.146099	0.000000	225.527364	
2028-01-01	19.502748	0.000000	216.884007	

✓ Saved: future_forecast.csv

```
In [31]: # Plot final forecast
fig, ax = plt.subplots(figsize=(16, 7))

# Historical data (last 120 months)
hist_plot = ts_data.iloc[-120:]
ax.plot(
    hist_plot.index,
    hist_plot.values,
    linewidth=1.5,
    color="black",
    label="Historical Data",
    alpha=0.8,
)

# Forecast
ax.plot(
    forecast_df["date"],
    forecast_df["forecast"],
    linewidth=2.5,
    color="red",
)
```

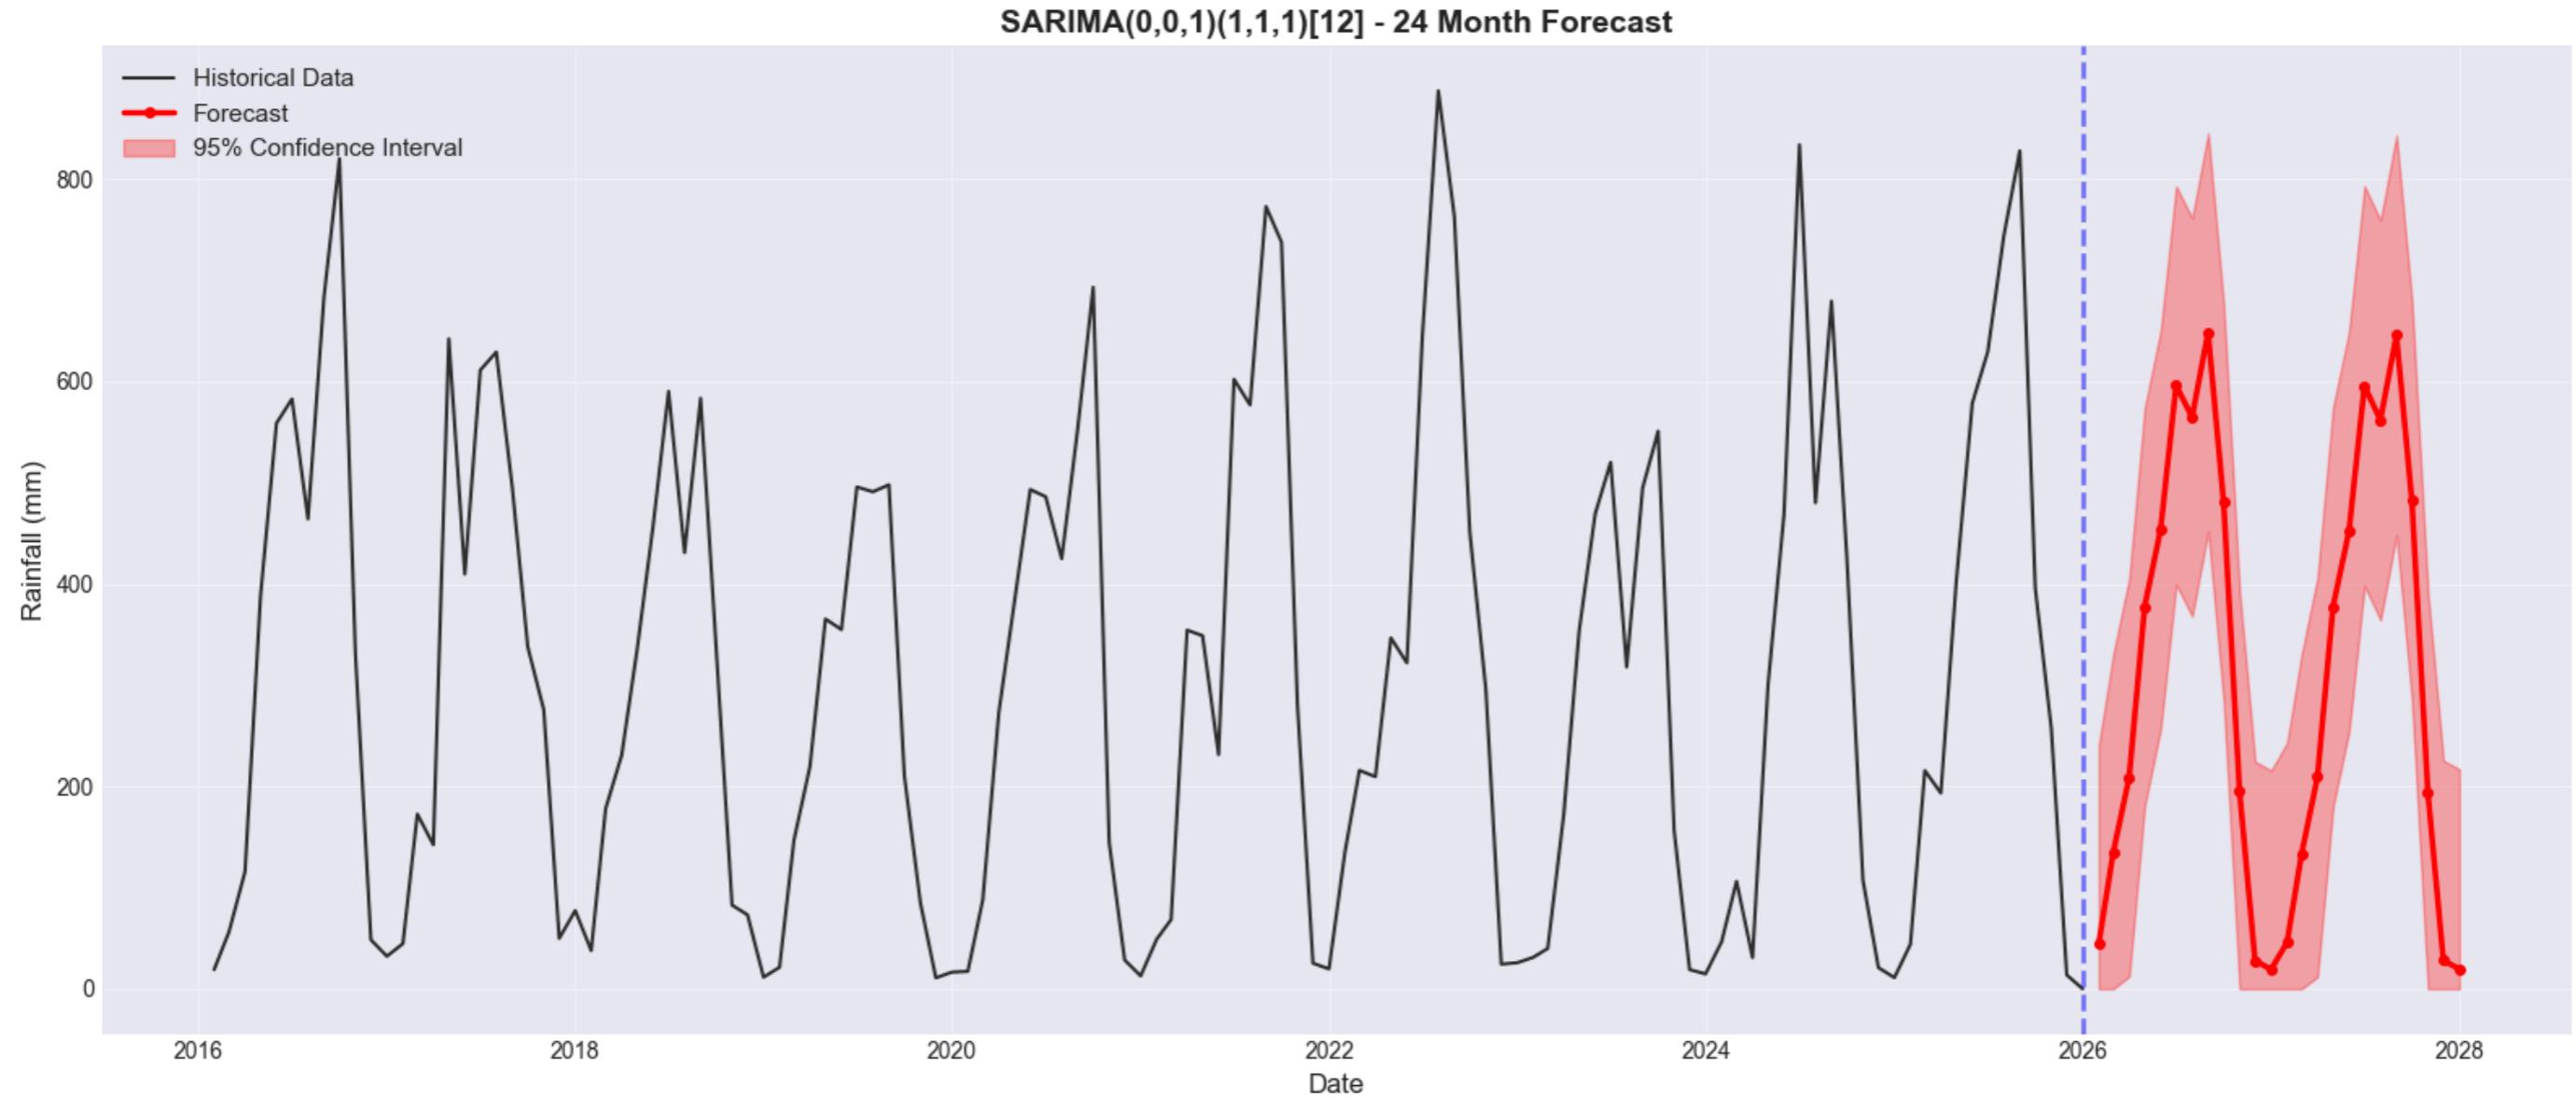
```
label="Forecast",
marker="o",
markersize=4,
)

# Confidence interval
ax.fill_between(
    forecast_df["date"],
    forecast_df["lower_ci"],
    forecast_df["upper_ci"],
    alpha=0.3,
    color="red",
    label="95% Confidence Interval",
)

# Vertical line at forecast start
ax.axvline(x=ts_data.index[-1], color="blue", linestyle="--", linewidth=2, alpha=0.5)

ax.set_title(
    f"{best_name} - {forecast_horizon} Month Forecast", fontsize=14, fontweight="bold"
)
ax.set_xlabel("Date", fontsize=12)
ax.set_ylabel("Rainfall (mm)", fontsize=12)
ax.legend(fontsize=11, loc="upper left")
ax.grid(True, alpha=0.3)

plt.tight_layout()
plt.savefig("final_forecast.png", dpi=300, bbox_inches="tight")
plt.show()
print("✓ Saved: final_forecast.png")
```



✓ Saved: final_forecast.png

14. Conclusions

Key Findings

1. Best Model Performance:

- The best performing model was identified through rigorous out-of-sample evaluation
- All models were evaluated using rolling one-step-ahead forecasts (no data leakage)

2. Seasonal Patterns:

- Strong annual seasonality (period = 12 months) was detected
- Seasonal differencing significantly improved stationarity
- SARIMA models generally outperformed non-seasonal models

3. Model Diagnostics:

- Residual analysis confirmed model adequacy
- Ljung-Box test validated absence of autocorrelation in residuals

- Forecast intervals provide uncertainty quantification

Methodological Improvements

✓ What was corrected:

- Removed deprecated `disp` parameter (statsmodels 0.14+ compatibility)
- Implemented proper train-test split (temporal separation)
- Used out-of-sample rolling forecasts (no data leakage)
- Applied seasonal differencing before modeling
- Conducted comprehensive residual diagnostics
- Compared models using proper evaluation metrics

Future Work

- Consider additional exogenous variables (temperature, pressure, etc.)
- Explore machine learning methods (LSTM, XGBoost) for comparison
- Implement automatic model selection (`auto_arima`)
- Extend forecast horizon with prediction interval analysis

✓ Analysis Complete

All outputs saved:

- `timeseries_plot.png`
- `decomposition.png`
- `seasonal_differencing.png`
- `acf_pacf.png`
- `train_test_split.png`
- `model_comparison.png` / `model_comparison.csv`
- `best_model_forecast.png`
- `residual_diagnostics.png`
- `final_forecast.png` / `future_forecast.csv`